Oblivious Linear-Time Tape Transposers Need Macroscopic Crossings

Hans Brende Claude Opus 5.5 (Anthropic) Gemini 4 Argon (Google DeepMind)

November 2026

Abstract

Can a multitape Turing machine transpose a $d \times d$ matrix of $\lceil \log_2 d \rceil$-bit entries in time linear in its length L? We show that an oblivious machine that does so must route a linear amount of information through *macroscopic crossings*: pairs of revisits of one tape, each taking the head a fixed fraction of L away from its cell, that interleave as a queue's would. Precisely, for every set of at least αL output bits, every set of steps that separates their output steps from the earlier ends of such crossings in the time–space dependency graph has at least c times as many elements as the output set. Here c and the fraction depend only on the number of tapes, the alphabet, the number of states, the time constant and α. The expansion can be found on a single coarsely monotone leg of a single tape, in a coarse skeleton of the run drawn from a finite family independent of d, and at least two work heads must make revisits of linear excursion. Behind these statements is a dichotomy for arbitrary linear-time traces: either such crossings are reached with linear expansion, or a separator exists that is cheaper than the acyclic set it leaves in the cone relation, which no transposer can have. For adaptive machines the conclusions hold on every trace followed by an exponentially small but fixed fraction of the inputs, and on the heavy traces of a bit simulator, with the same tapes and alphabet, of every linear-time transposer; carrying them back to the machine's own traces is open. We also prove that every cut between a set of input rows and the outputs they feed into a set of output rows carries the corresponding information whatever the coding, and we exclude classes of machines with boundedly many macroscopic phases, oblivious and adaptive. Explicit linear-time schedules, realized by machines that compute other functions, meet these conditions in every test we could run (numerically, at $d \leq 128$), so the conditions are not known to decide the question. We show that one set of steps separates the inputs from the outputs for every pair of row sets at once, so counting such cuts cannot by itself go beyond linear time, and we state the multi-scale question that remains. The main theorems are proved in Lean 4; one proposition and one corollary are proved in the text, and the exclusions of individual schedule families are given as proof sketches. The constants of the main theorems are explicit and astronomically small; those of the two-heads theorem are polynomial, and those of one class exclusion come from a compactness argument.

Contributions. Hans Brende chose the problem, fixed the verification standard, and directed the research. Gemini 4 Argon, a large language model developed by Google DeepMind, developed most of the mathematics and wrote the Lean formalization. Claude Opus 5.5, a large language model developed by Anthropic, reviewed the mathematics and the Lean throughout, refereed the main proofs, proposed several of the statements proved here (among them the macroscopic form of the crossing theorem and the trace-only dichotomy), and wrote this paper under Hans Brende's direction.

Contents

1 Introduction **4**

1.1 The problem . . . 4
1.2 Results . . . 4
1.3 How to read the paper and its formal proofs . . . 6
1.4 Related work . . . 6
1.5 Organization . . . 7

2 The model, traces and the dependency DAG **7**
2.1 Machines and transposition . . . 7
2.2 Traces . . . 8
2.3 Revisits, excursions and crossings . . . 8
2.4 The dependency DAG, cuts and cone relations . . . 9

3 The certificate method **10**
3.1 Genies and index codes . . . 10
3.2 Cruxes . . . 10
3.3 Heavy traces of arbitrary machines . . . 11
3.4 Tools for finding acyclic sets . . . 12

4 Two long-range heads **13**

5 Macroscopic crossings **15**
5.1 Constants . . . 15
5.2 The time-pullback bound . . . 15
5.3 The bridge . . . 16
5.4 Theorem D: a dichotomy for every linear-time trace . . . 17
5.5 Theorem X . . . 17

6 Skeletons: where the crossings must be **18**
6.1 Crossings are zigzags . . . 19
6.2 Skeletons . . . 19
6.3 Theorem S . . . 21

7 Rows to columns **22**

8 Adaptive machines **23**

9 Class exclusions **25**
9.1 Boundedly many macroscopic phases . . . 25
9.2 A native gap suffices . . . 26
9.3 Few scales and no free head . . . 27

10 What a linear-time transposer needs **27**
10.1 The specification . . . 27
10.2 Mechanisms that fail . . . 28
10.3 Schedules that meet every condition we could test . . . 28
10.4 Why cuts alone cannot finish . . . 29

11 Open problems **30**

A Catalogue of excluded schedule families **31**

B Index of Lean declarations **34**

C Building and verifying the formalization **34**

1 Introduction

1.1 The problem

Can a deterministic multitape Turing machine transpose a $d \times d$ matrix with $b(d) = \max(1, \lceil \log_2 d \rceil)$-bit entries, written row by row, in time linear in the input length $L = d^2 b(d)$? This is the statement (T) of the first paper of this series [9]. By a reduction of Harvey and van der Hoeven [21, Cor. 6.9], a linear-time multiplication machine would give a linear-time transposer, so a proof of $\neg$(T) would give the first superlinear lower bound for off-line integer multiplication on multitape Turing machines. The question is open, and this paper does not decide it.

A machine is *oblivious* if its head motion and the steps at which it appends output bits depend only on d. Every machine can be made oblivious at a logarithmic cost in time [28], so oblivious machines are the natural first target. The run of an oblivious machine is a fixed time–space network, and a linear-time oblivious transposer is a network code for the transposition demands on that network. Adler, Harvey, Jain, Kleinberg and Lehman showed that if coding could not beat routing on undirected networks, oblivious two-tape transposition would need time $\Omega(L \log d)$ [1, §6]; Braverman and He, and concurrently Zhang, Li and Li, have refuted that hypothesis [4, 30], so for oblivious machines the question is genuinely one about codes. The series has reduced it to codes in two steps. The fourth paper reduced $\neg$(T) for oblivious machines to depth reduction and a coding statement, the *crux* [7, Theorem D]. Lower bounds for several restricted classes of labels suggested that depth reduction might fail [7, 5, 12, 10], but the fourteenth paper proved it for every number of heads, so that the crux alone would exclude oblivious linear-time transposers [6, Corollary O]. The crux asks whether the *cone relations* of a linear-time schedule, the relations between each output bit and the input bits it can see without crossing a small revealed set, admit codes of rate close to 1. Machine-free relations of this kind carry storage codes above rate one half [11], and binary codes of rate approaching one exist on suitable separated relations [14]; whether the relations of an actual linear-time schedule can carry them is the question.

This paper attacks the crux from the side of the schedule. Instead of a coding bound for all relations, we ask which schedules make the crux *provable*. We show that it holds, with an explicit deficit, for every schedule in which a small enough set of steps separates a large set of outputs from the *macroscopic crossings* of the tapes. An oblivious linear-time transposer must therefore route a linear amount of information through macroscopic crossings, and we describe where it must do so.

1.2 Results

The run of a machine defines a *dependency DAG* on its steps $0, \dots, T-1$: every step depends on the step before it, which carries the state, and on the previous visit of every work head to the cell it is scanning, which carries that cell's contents. A *revisit* of a work tape is a pair $u < t$ of consecutive visits of its head to one cell between which the head leaves the cell; its *excursion* is the farthest the head travels from the cell before returning. Two revisits $u < t$ and $u' < t'$ of the same tape *cross* if $u < u' < t < t'$. A unit-speed head makes crossing revisits only when it reuses a stretch of tape like a queue, across two turns: write a stretch moving right, return, and read it again moving right. Reuse like a stack, across one turn, makes only nested revisits. Given $\theta > 0$, a step is a *macroscopic crossing step* if it is the earlier end of a revisit of excursion more than θL that crosses another revisit of excursion more than θL on the same tape. A trace of a machine is η-*heavy* if at least a $2^{-\eta L}$ fraction of the inputs follow it; an oblivious machine has one trace, and it is 0-heavy. Section 2 states these definitions exactly.

- **Theorem X, macroscopic crossings** (Theorem 5.5). Let M be an oblivious machine with

k work tapes that transposes $d \times d$ matrices within $C(L+1)$ steps, and let $\alpha > 0$. There are explicit $\theta > 0$, $c > 0$ and d_0, depending only on k, the alphabet and state counts of M, C and α, such that for every $d \geq d_0$ the following holds. For every set V_0 of at least αL output positions, every set of steps that meets every dependency path from the output steps of V_0 to the macroscopic crossing steps has at least $c|V_0|$ elements. By Menger's theorem, there are then $c|V_0|$ vertex-disjoint dependency paths from these outputs into macroscopic crossings.

- **Theorem D, a dichotomy for all traces** (Theorem 5.4). Theorem X is the transposer case of a statement about arbitrary linear-time traces, with no hypothesis on what the machine computes: for every output set of size at least αL, either the macroscopic crossings are reached with linear expansion, or a cheap separator leaves a linear-size acyclic set in the cone relation. A separator cheaper than the acyclic set it leaves, at a radius at least the genie depth, is impossible for a transposer (Lemma 3.5).

- **Theorem S, skeletons** (Theorem 6.7). At scale $\theta L/4$ the walk of each work tape of a linear-time run is a word of at most $M_{\max} = \lfloor 16C/\theta \rfloor$ steps ± 1, so the coarse shapes of all runs come from an explicit finite family of $(2^{M_{\max}+1} - 1)^k$ tuples, independent of d. For an oblivious linear-time transposer the shape is *crossing-capable* (some tape turns at least twice at this scale with overlapping legs), and for every large output set the expansion of Theorem X is carried by the crossing steps of a *single coarsely monotone leg of a single tape*. Along infinitely many sizes, one fixed shape, tape and leg serve for the set of all input positions.

- **Theorem H, two long-range heads** (Theorem 4.1). No machine that transposes in linear time with an η-heavy trace, for a small fixed η, has on that trace one unrestricted work head while every other work head keeps its revisits within a small fixed fraction of L, whatever the input head does. Every such transposer, in particular every oblivious one, therefore has at least two work heads that return to cells after moving a fixed fraction of L away; and for each work head, some set containing a fixed fraction of its outputs reaches long revisits off that head with linear expansion.

- **Theorem E, adaptive machines** (Theorem 8.1). The conclusion of Theorem X holds for every η-heavy trace of an arbitrary linear-time transposer, for an explicit $\eta > 0$. It holds on all heavy traces of the bit simulator of the fifteenth paper [14, Theorem S] applied to any linear-time transposer; that simulator has the same tapes and work alphabet, runs in linear time, and its heavy traces carry all but a $2^{-\eta L}$ fraction of the inputs. Transferring the conclusion back to the original machine's own dependency DAG is open; every statement in this paper about the dependency DAG of a general machine assumes a heavy trace or names the simulator.

- **Lemma R, rows to columns** (Lemma 7.1). For every machine that transposes at size d, every η-heavy trace, and every two sets A_0, A_1 of rows, every set of steps that separates the input reads of the rows in A_0 from the output steps of the entries that go from rows A_0 to output rows A_1 has at least $(|A_0||A_1|b - \eta L)/(w_{\text{val}} + 1)$ elements, where w_{val} is the number of bits of one step's value.

- **Class exclusions** (Section 9). Machines whose heads make boundedly many macroscopic phases, oblivious or adaptive with sublinear jitter, are excluded by an explicit native deficit (Theorems B and J); a fixed native deficit suffices in place of the half-rate hypothesis of the fourteenth paper's low-adaptivity corollary (Theorem G); and machines with no free work head, bounded nesting load and boundedly many scales are excluded by a scale-gap pigeonhole (Theorem Q).

Appendix A lists explicit schedule families excluded by these theorems and by certificates, with proof sketches.

What remains. Explicit oblivious linear-time schedules, realized in our experiments by machines that compute other functions, meet the necessary conditions above in every test we could run (numerically, at $d \leq 128$; Section 10). So these conditions are not known to decide the question. Counting cuts alone cannot go further: one set of steps, the storage live across the last input read together with the earlier output steps, separates the inputs from the outputs for every pair of row sets at once (Proposition 10.2). Outside restricted schedule classes (Corollary 7.3), the only coding-proof logarithmic factor known on tapes comes from counting passes (Theorem P of the first paper), and Section 11 states the multi-scale question it leaves open.

Constants. The constants of Theorems X, D and S are explicit, positive and astronomically small. For two work tapes, time constant $C = 1$ and $\alpha = 1$, the expansion constant of Theorem X is $\rho_0/(4w_{\mathrm{val}})$ with $\rho_0 \approx 10^{-105}$, and the threshold θ is smaller still; Section 5 gives the formulas. The constants of Theorem H are polynomial in the time constant and the value width, and those of Theorem Q come from a compactness argument and are not explicit. The small constants come from the proofs: on the schedules we measured, the expansion ratio was exactly 1 at every tested threshold up to $L/2$ (numerically, at $d \leq 128$; Remark 5.8).

1.3 How to read the paper and its formal proofs

Every numbered result states its status. Lean: lines name the declarations of the accompanying Lean 4 library that prove it, in the file named at the right; the Lean statement is authoritative and the text gives a faithful outline of the formal proof. A result marked Not formalized has a complete written proof in the text. Cited: lines name results of earlier papers of the series used without proof, and say whether they are Lean-verified there. Numerical observations appear only in remarks or parentheses labelled as numerical and are never used in proofs. The formal library depends only on the axioms `propext`, `Classical.choice` and `Quot.sound` (Appendix C).

The formalization is in the frozen machine model of the first paper [9], through the traces, dependency DAGs and index codes of the eleventh paper [10], the genies of the fourteenth [6] and the simulator of the fifteenth [14]. Not formalized: Proposition 10.2 and Corollary 7.3, which have complete proofs in the text; the remarks marked as proof sketches; and the exclusions of the individual schedule families of Appendix A, which are not numbered results and are given as proof sketches over the formal class theorems and certificate lemmas.

1.4 Related work

Lower bounds on tapes. Superlinear lower bounds for explicit problems are known for one work tape: Hennie's crossing sequences [22], Maass's combinatorial arguments [25], and for transposition itself the bounds of Dietzfelbinger, Maass and Schnitger [17] and Dietzfelbinger and Maass [16] in the classical one-tape models. With two or more work tapes the only superlinear lower bounds known are those that follow from the separation of deterministic and nondeterministic linear time [27], for problems that are hard for nondeterministic linear time [18]; none is known for off-line multiplication, sorting, the discrete Fourier transform or transposition [21, §1.4]. The simulation of time by space [23] also rests on the structure of multitape computation graphs. For transposition, moving the entries as indivisible items along a fixed pattern needs time $\Omega(L \log d)$ [29], and similar bounds hold for indivisible records in external memory [2]; Kirchherr proved an $\Omega(\log d)$ bound

on reversals for conservative machines [24], and the first paper of this series an $\Omega(\log d)$ bound on monotone passes for arbitrary computation, with coding allowed [9, Theorems P and P′], extended to adaptive machines without a runtime assumption in [13]. The stream lower bounds of Grohe, Hernich and Schweikardt [20, 19] bound reversals for sorting-type problems with small internal memory. A linear-time machine, however, may reverse its heads $\Theta(L)$ times at small scales, so pass and reversal bounds apply to it only after its small-scale motion is absorbed into the finite control.

Macroscopic motion. The first paper's Theorem M* [9, Theorems 7.8 and 7.9] passes to the limit of rescaled head trajectories and shows that the heads of any linear-time transposer, oblivious or adaptive, satisfy an integral inequality forcing returns to macroscopically earlier positions. The fourth paper deduced, by a written proof, that in every macroscopic limit of a linear-time transposer at least two work-head limits are non-monotone [7, Corollaries 3.18 and 3.24]. Theorem H(b) is a quantitative form of this at a single size, with explicit constants and on heavy traces. Theorem X asks for more than non-monotone motion: a run whose heads reuse their tapes only like stacks has no macroscopic crossing steps, so Theorem X excludes it whatever its macroscopic limit.

Network and index coding. The time–space view of oblivious machines and the conditional lower bound are due to Adler et al. [1]; the multiple-unicast conjecture was refuted by Braverman and He [4] and, concurrently, by Zhang, Li and Li [30]. Our exclusions use the acyclic-set lower bound of index coding [3], in the machine form of the eleventh paper [10].

Upper bounds. Recursive block transposition takes time $O(L \log d)$. An unrefereed preprint claims time $O(L(\log d)^{1-\kappa})$ for a small $\kappa > 0$ through an address-interchange primitive that mixes data by exclusive or [26]; if it is correct, coding already beats routing on tapes by a factor $(\log d)^{\kappa}$, and any proof of $\neg$(T) must allow such gains.

1.5 Organization

Section 2 sets out the model, the dependency DAG, revisits and crossings, cuts and cone relations. Section 3 gives the certificate method: an acyclic set of linear size in a cone relation, possibly after a cheap cut, excludes linear-time transposition. Section 4 proves Theorem H, Section 5 Theorems D and X, and Section 6 Theorem S. Section 7 proves Lemma R and the colouring criterion built on it. Section 8 treats adaptive machines, Section 9 the class exclusions. Section 10 turns the results around into a specification for a linear-time transposer and describes the schedules that meet it, and Section 11 lists the open problems that would decide the question. Appendix A catalogues the excluded schedule families, Appendix B indexes the Lean declarations, and Appendix C explains how to build and check the formalization.

2 The model, traces and the dependency DAG

2.1 Machines and transposition

We work in the frozen machine model of the first paper [9]. A machine M has a read-only two-way input tape, $k = k(M)$ work tapes over the alphabet $\{0, \dots, g\}$ with blank 0, states $\{0, \dots, q\}$ with start state 0, and a one-way write-only output tape that receives at most one bit per step. Every step writes a symbol on every work tape (possibly the scanned one) and moves every head by at most one cell. The *value width*

$$w_{\mathrm{val}}(M) = \lceil \log_2(q+1) \rceil + k \lceil \log_2(g+1) \rceil + 2$$

bounds the number of bits needed to record what one step does: the new state, the k written symbols and the appended output bit. Always $w_{\mathrm{val}} \geq 2$.

For a side length $d \geq 1$ put $b = b(d) = \max(1, \lceil \log_2 d \rceil)$, $W = db$ and $L = d^2 b$. A $d \times d$ matrix with b-bit entries is written row by row as a word of L bits: the bit γ of entry (i, j) sits at position $(di + j)b + \gamma$. The transpose has the same encoding with i and j exchanged, so output position $\mathrm{tpos}(q)$ carries input bit q, where

$$\mathrm{tpos}\big((di + j)b + \gamma\big) = (dj + i)b + \gamma,$$

an involution of $[0, L)$. For $q = (di + j)b + \gamma$ we call i the *input row* of q and j its *output row*: the copy of q is written in row j of the output. *M transposes at size d within time c*, written $\mathrm{TransposesAt}(M, d, c)$, if on every matrix input it halts within $\lfloor c(L + 1) \rfloor$ steps with the encoding of the transpose on its output tape. Statement (T) asks for a machine and a constant c with this property at every d.

M is *oblivious at length L with halting time T* if every input of length L halts at exactly step T and any two such inputs have the same input-head position, the same work-head positions and the same append decision at every step. We write $\mathrm{ObliviousLinTime}(M, d, C, T)$ if M is oblivious at length L with halting time $T \leq C(L + 1)$ and transposes at size d. This is a condition at a single size d.

2.2 Traces

The *trace* of M on an input w records, step by step, the move of the input head, the moves of the work heads and whether a bit is appended [10]. It determines the positions $y_j(t)$ of work head j after t steps, with $y_j(0) = 0$ and $|y_j(t + 1) - y_j(t)| \leq 1$, the input-head position, and the *append steps* $u_0 < u_1 < \dots$ at which output bits are written. For an oblivious machine all inputs of length L share one trace, the trace on the all-zero input.

For a trace τ the *fibre* of τ is the set of inputs $X \in \{0, 1\}^L$ whose trace is τ, and the *trace probability* $\Pr(\tau)$ is the size of the fibre divided by 2^L. A trace is *η-heavy* if $\Pr(\tau) \geq 2^{-\eta L}$. An oblivious machine has a single trace, of probability 1.

If a run appends exactly L bits before step T, let $a(r)$ be the step at which output bit r is appended. For an input position q the *root* of q is

$$u_0(q) = a(\mathrm{tpos}(q)),$$

the step at which the output copy of input bit q is written, and $u_0(V) = \{u_0(q) : q \in V\}$ for $V \subseteq [0, L)$. The map u_0 is injective, because append steps are distinct and tpos is a bijection (`appendTime_tposOf_injective`). Results quantified over sets V_0 of input positions are therefore statements about the output steps of the corresponding output sets $\mathrm{tpos}(V_0)$.

2.3 Revisits, excursions and crossings

Fix a work tape, with head positions $y(t)$, and a horizon T.

Definition 2.1 (Revisits)**.** Lean: `BoundedNesting.IsHop`, `HasHopExcursionGTAt`, `HopsCross` [`EnvelopeCrux.lean`]

A *revisit* (in Lean, a hop) is a pair of steps $p < u < T$ with $p + 2 \leq u$, $y(u) = y(p)$, and $y(t) \neq y(p)$ for all $p < t < u$: two consecutive visits to one cell between which the head leaves it. Its *excursion* is $\max_{p<t<u} |y(t) - y(p)|$, the farthest the head travels from the cell before returning. Its *earlier end* is p. Two revisits (p, u) and (p', u') of the same tape *cross* if $p < p' < u < u'$ or $p' < p < u' < u$.

A revisit lies on one side of its cell: the head stays strictly above it, or strictly below it, in between. Two crossing revisits lie on opposite sides (Lemma 6.1(a); compare [5, §2]). Every step is the earlier end of at most one revisit on each tape.

Definition 2.2 (Crossing steps)**.** Lean: `crossingLowerEtaSteps` [`ConjectureEBridge.lean`]
For a trace with horizon T and an integer $H \geq 1$, let CL_H be the set of steps $v < T$ such that, on some work tape, v is the earlier end of a revisit of excursion greater than $2H$ that crosses another revisit of excursion greater than $2H$ on the same tape.

In Lean CL_H is `crossingLowerEtaSteps` $\tau\, T\, H\, 1\, T$. The set shrinks as H grows (`crossingLowerEtaSteps_anti_H`), so a lower bound on the cost of separating from CL_H is stronger for larger H.

To see what CL_H measures, consider a head that sweeps right across a stretch of tape, turns, sweeps back, turns again and sweeps right. A cell visited on the first sweep and revisited on the second is the earlier end of a revisit above it; a cell visited on the second sweep and revisited on the third is the earlier end of a revisit below it; and when the second turn overlaps the first stretch these revisits cross. This is how a tape is used as a queue. A tape used as a stack, written while the head moves one way and read back while it returns, makes only nested revisits. Section 6 turns this picture into theorems.

2.4 The dependency DAG, cuts and cone relations

Definition 2.3 (Dependency DAG)**.** Lean: `MergeDepth.DepStep`, `MergeDepth.IsDepPath` [`MergeDepth/ColdGenie.lean`]
The *dependency DAG* of a trace with horizon T has the vertices $0, \ldots, T-1$ and two kinds of edges: the *state edge* $t \to t-1$ for $1 \leq t < T$, and the *revisit edge* $u \to p$ for every revisit (p, u) of every work tape. A *dependency path* is a sequence $t_0 \to t_1 \to \cdots \to t_m$ of edges.

Every edge decreases the step, so the DAG is acyclic. The value written at step t is a function of the values at its out-neighbours and of the input bit read at t: the state edge carries the state, and the revisit edges carry the contents of the scanned cells. In Lean the edge relation is `DepStep` and paths are `IsDepPath` (from the eleventh paper's library [10]).

Definition 2.4 (Cuts)**.** Lean: `IsMengerVertexCut` [`OneFreeHeadSmallExcursion.lean`]
For sets S and D of steps, a set B of steps *separates S from D* if every dependency path that starts in S and ends in D contains a vertex of B, its endpoints included. In Lean this is `IsMengerVertexCut`.

By Menger's theorem the least size of a set separating S from D equals the largest number of vertex-disjoint dependency paths from S to D. We use Menger's theorem only to explain statements; the theorems and their formal proofs bound the size of separating sets directly.

Definition 2.5 (Cones, balls and acyclic sets)**.** Lean: `BoundedStackDepth.depCone`, `BoundedStackDepth.coneSystem`, `ballAvoidSystem`, `HalfRate.IsAcyclicIn` [`EnvelopeCrux.lean`]
Let Γ be a set of steps. For a step t, the *cone* of t avoiding Γ is the set of input positions p such that some dependency path $t \to t_1 \to \cdots \to t_m$ with $t_1, \ldots, t_m \notin \Gamma$ ends at a step t_m (or at t itself, if $m = 0$) at which the input head reads position p. The root t itself need not avoid Γ. The *cone relation* assigns to each input position q the cone of its root $u_0(q)$ (in Lean, `coneSystem` evaluated at $\mathrm{tpos}(q)$). The *ball relation of radius D avoiding B* is defined in the same way using only paths with at most D edges after the root (`ballAvoidSystem`). A set U of input positions is *acyclic* in a relation N if every nonempty $S \subseteq U$ has an element v with $N(v) \cap S = \emptyset$ (`HalfRate.IsAcyclicIn`).

An acyclic set is a set of receivers none of which can see, in its cone, the bits wanted by the receivers that come after it in a suitable order. In index coding, a message that lets every receiver of an acyclic set recover its bit must carry at least as many bits as the set has elements [3]. For machines this takes the following form: what the cut steps write, together with what a small *genie* reveals, is a message from which each output of the acyclic set is computed using only its cone. The fourteenth paper's genie (Lemma 3.1 below) bounds the depth of the cones at small cost, and the eleventh paper's index-code lemma turns an oblivious transposer into such a message [10].

3 The certificate method

Every exclusion in this paper has the same shape: find, in the cone relation of a run, a linear-size acyclic set, possibly after revealing a cheap set of steps, and conclude that the run cannot belong to a linear-time transposer. This section states the method and the tools used to find acyclic sets.

3.1 Genies and index codes

Lemma 3.1 (Genies and index codes)**.** Cited: [6, Corollary O], Lean-verified there (`BoundedStackDepth.genie_cost`, `BoundedStackDepth.genie_of_trace`, `BoundedStackDepth.exists_indexCode`).

(a) For every $k \geq 1$, $\beta \geq 1$, $C \geq 1$, $\varepsilon > 0$ and every trace with k work tapes and $T \leq C(L+1)$, there is a set Γ of steps with $\beta|\Gamma| \leq \varepsilon L$ such that every dependency path whose vertices after the first avoid Γ has at most $\mathrm{genieDepth}(k, \beta, C, \varepsilon)$ *vertices. The depth bound depends only on k, β, C and ε.*

(b) If M is oblivious at length L with common trace τ and transposes, then for every set Γ of steps the values written at the steps of Γ form an index code of $w_{\mathrm{val}}|\Gamma|$ bits for the cone relation avoiding Γ: two inputs with the same values on Γ that agree on the cone of $u_0(q)$ agree on the output bit at $u_0(q)$.

The genie comes from the fourteenth paper's solution of the trajectory cut problem [6, Theorem T]; its depth $\mathrm{genieDepth}(k, \beta, C, \varepsilon)$ is explicit, with $\log\log \mathrm{genieDepth} = O(k^2 \log(k\beta C/\varepsilon))$ [6, Remark 9.4 and §10].

Lemma 3.2 (Acyclic sets bound codes)**.** Lean: `storageCode_card_le_of_isAcyclicIn`, `indexCode_bits_ge_of_isAcyclicIn` [`DichromaticCrux.lean`]

Let N be a relation on a finite set V and $A \subseteq V$ acyclic in N. Every storage code for N over an alphabet Σ has at most $|\Sigma|^{|V|-|A|}$ words, and every index code for N and the demand tpos *has at least $|A|$ bits, provided A is acyclic for $q \mapsto N(\mathrm{tpos}(q))$.*

Outline of the formal proof. On a storage code a word is determined by its letters outside A: recover the letters in A in an order in which each letter's neighbourhood lies outside the not-yet-recovered part (`acyclicIn_outside_injOn`). The fibres of an index code are storage codes for the demand-shifted relation, and pigeonhole over the messages gives the bound. □

3.2 Cruxes

Definition 3.3 (Cruxes)**.** Lean: `MachineCruxAt`, `MachineCutCruxAt`, `SubpopCutCruxAt` [`EnvelopeCrux.lean`]

Let M be a machine, $C \geq 1$, $D \geq 0$ and $\rho, \delta, \alpha_0 > 0$.

(a) M has *crux rate ρ at depth D* if for every $\eta > 0$ and all large d, whenever M is oblivious at length L with halting time $T \le C(L+1)$ and exactly L appends, every index code for the cone relation avoiding any Γ that cuts all dependency paths to at most $D+1$ vertices has at least $(\rho-\eta)L$ bits (`MachineCruxAt`).

(b) M has *cut-crux rate ρ at cost δ and depth D* if, under the same conditions, there is a set B of steps with $w_{\mathrm{val}}|B| \le \delta L$ such that every index code for the cone relation avoiding $\Gamma \cup B$ has at least $(\rho-\eta)L$ bits (`MachineCutCruxAt`).

(c) M has *subpopulation cut-crux rate ρ at density α_0, cost δ and depth D* if, under the same conditions, there are a set V_0 of at least $\alpha_0 L$ input positions and a set B with $w_{\mathrm{val}}|B| \le \delta|V_0|$ such that every index code for the cone relation avoiding $\Gamma \cup B$ has at least $(\rho-\eta)|V_0|$ bits (`SubpopCutCruxAt`).

Theorem 3.4 (Cruxes exclude)**.** Lean: `subpop_cutcrux_excludes`, `machine_cutcrux_excludes` [`EnvelopeCrux.lean`]

Let $C \ge 1$ and $\varepsilon > 0$, and put $D_\varepsilon = \mathrm{genieDepth}(k, w_{\mathrm{val}}, C, \varepsilon)$. If M has crux rate $\rho > \varepsilon$ at depth D_ε, or cut-crux rate ρ at cost δ with $\varepsilon + \delta < \rho$, or subpopulation cut-crux rate ρ at density α_0 and cost δ with $\varepsilon < \alpha_0(\rho - \delta)$, then for all large d there is no T with $\mathrm{ObliviousLinTime}(M, d, C, T)$.

Outline of the formal proof. Suppose M transposes obliviously in time $C(L+1)$. Reveal the genie Γ of Lemma 3.1(a), of cost εL and depth D_ε, together with B. By Lemma 3.1(b) the values on $\Gamma \cup B$ form an index code of at most $(\varepsilon+\delta)L$ bits, respectively $\varepsilon L + \delta|V_0|$ bits, contradicting the crux for small η. The crux form is `machine_crux_excludes` (`MachineCrux.lean`). □

Acyclic sets give cruxes. If, for every $\eta > 0$ and all large d, every cone relation of depth D in Definition 3.3(a) contains an acyclic set of at least $(\rho-\eta)L$ input positions, then M has crux rate ρ at depth D, by Lemma 3.2 (`machineCrux_of_isAcyclicIn`, `oblivious_excludes_of_isAcyclicIn`). The same holds for ball relations, which contain the cone relations of depth D (`machineCutCruxAt_of_ballAvoid_isAcyclicIn`).

The next lemma is the form used by Theorems X and S: at a single size, revealing a set B that leaves an acyclic set U in the ball relation costs at least $|U|$ bits, up to the genie.

Lemma 3.5 (Acyclic sets cost cuts)**.** Lean: `obliviousLinTime_separator_card_ge_of_isAcyclicIn`, `obliviousLinTime_ballSeparator_card_ge` [`EnvelopeCrux.lean`]

Let $\mathrm{ObliviousLinTime}(M, d, C, T)$ *with $d \ge 1$, $C \ge 1$ and $\varepsilon > 0$. For every set B of steps and every set U of input positions that is acyclic in the ball relation of radius* $\mathrm{genieDepth}(k, w_{\mathrm{val}}, C, \varepsilon)$ *avoiding B,*

$$w_{\mathrm{val}}|B| \;\ge\; |U| - \varepsilon L.$$

Outline of the formal proof. Reveal the genie of cost εL together with B. The cone relation avoiding both is contained in the ball relation (`coneSystem_subset_ballAvoidSystem`), so U is acyclic in it, and Lemma 3.1(b) with Lemma 3.2 gives $w_{\mathrm{val}}(|\Gamma| + |B|) \ge |U|$. □

3.3 Heavy traces of arbitrary machines

For a machine that is not oblivious, the index code of Lemma 3.1(b) is not available, but counting the fibre of a single trace replaces it.

Lemma 3.6 (Frozen-trace deficit)**.** Lean: `frozenTrace_isAcyclicIn_coneSystem_deficit_le`, `frozenTrace_isAcyclicIn_ballAvoid_separator_card_ge` [`SimulatorBridge.lean`]

Let M transpose at size $d \geq 1$ within time c, and let τ be an η-heavy trace of M with halting time T. Then $T \leq 2cL$, τ appends exactly L bits, and for every set Γ of steps and every set U acyclic in the cone relation of τ avoiding Γ,

$$|U| - w_{\mathrm{val}}|\Gamma| \;\leq\; \eta L.$$

Outline of the formal proof. On the fibre of τ all inputs share the dependency DAG. Split the fibre according to the $2^{w_{\mathrm{val}}|\Gamma|}$ possible values on Γ. Each part is a storage code for the cone relation, so it has at most $2^{L-|U|}$ elements by Lemma 3.2. Hence $2^{(1-\eta)L} \leq |\text{fibre}| \leq 2^{w_{\mathrm{val}}|\Gamma|+L-|U|}$. □

3.4 Tools for finding acyclic sets

Acyclic sets in cone relations are found by potentials. The following statements hold for any relation N on a finite set; an arc $q \to p$ means $p \in N(q)$.

Lemma 3.7 (Envelopes and windows)**.** Lean: `acyclicIn_of_lt_on_nbhd`, `upperEnvelope_isAcyclicIn`, `oneSidedEnvelope_isAcyclicIn`, `window_isAcyclicIn_of_noFarArcs`, `exists_isAcyclicIn_of_window_maxDeg` [`EnvelopeCrux.lean`]

(a) If a potential strictly decreases along every arc inside S, then S is acyclic.

(b) Let φ, ψ be real potentials and $w \geq 0$. The window $\{q \in V_0 : q \notin N(q),\ \varphi(q) + 2w < \psi(q)\}$ *is acyclic if every arc $q \to p$ inside it is* near, $|\psi(p) - \varphi(q)| \leq w$. *If instead every vertex of the window has at most β far arcs inside it, the window contains an acyclic set of at least a $1/(\beta+1)$ fraction of its elements.*

Outline of the formal proof. (a) Inside S an element of least potential has no out-arc into S. (b) A near arc $q \to p$ inside the window has $\varphi(p) < \psi(p) - 2w \leq \varphi(q) - w$, so φ strictly decreases along it and (a) applies. With far arcs, remove the far neighbours of the vertices one at a time; a greedy choice keeps a $1/(\beta+1)$ fraction. □

Lemma 3.8 (Cut sets and block ranking)**.** Lean: `cutSet_isAcyclicIn`, `exists_isAcyclicIn_card_ge_of_blockCut`, `exists_isAcyclicIn_card_ge_of_uniformBlockCut`, `exists_isAcyclicIn_of_shiftBlockCutFamily_gen` [`EnvelopeCrux.lean`]

Assign to each $q \in V_0$ two disjoint sets $A(q), E(q)$ of blocks in $\{0, \dots, B-1\}$, and suppose every arc $q \to p$ inside V_0 has $A(q) \cap E(p) \neq \emptyset$. For a ranking σ of the blocks, the set of q with $\sigma(A(q)) < \sigma(E(q))$ (every block of $A(q)$ ranked below every block of $E(q)$) is acyclic, and some ranking makes it contain at least $\sum_q \binom{|A(q)|+|E(q)|}{|A(q)|}^{-1}$ elements, hence at least $|V_0|/\binom{a+e}{a}$ elements when $|A(q)| \leq a$ and $|E(q)| \leq e$ throughout.

Outline of the formal proof. Along an arc $q \to p$ inside the set, a block of $A(q) \cap E(p)$ shows that the largest rank in $A(p)$ is below the smallest rank in $E(p)$, hence below that block's rank, hence at most the largest rank in $A(q)$; so the largest rank of A strictly decreases along arcs, and Lemma 3.7(a) applies. For disjoint $A(q), E(q)$, exactly a fraction $\binom{|A(q)|+|E(q)|}{|A(q)|}^{-1}$ of all rankings put $A(q)$ below $E(q)$ (`choose_mul_card_goodPerms_eq_factorial`), so the average size of the set over all rankings is the stated sum. □

Lemma 3.9 (Step-invariant windows)**.** Lean: `isAcyclicIn_ballAvoid_window_of_stepInvariant`, `isAcyclicIn_ballAvoid_window_of_constStepInvariant` [`EnvelopeCrux.lean`]

Let a trace, a set B of steps, a depth D and a set V_0 of input positions be given, with regions $R(0), \dots, R(D)$ of steps, a sink *region, a step potential μ and potentials φ, ψ on input positions. Suppose: the roots of V_0 lie in $R(D)$ with $|\mu(u_0(q)) - \varphi(q)| \le c_0$; every edge from the sink that avoids B stays in the sink, and the sink reads no position of V_0; every edge from $R(r+1)$ that avoids B enters the sink or enters $R(r)$ changing μ by at most Δ; and a read of $p \in V_0$ at a step $u \in R(r)$ has $|\psi(p) - \mu(u)| \le c_1$. Then $\{q \in V_0 : \varphi(q) + 2(c_0 + D\Delta + c_1) < \psi(q)\}$ is acyclic in the ball relation of radius D avoiding B.*

Outline of the formal proof. By induction along a dependency path from a root with at most D edges that avoids B after the root, its end lies in the sink or in the region $R(D-\ell)$ after ℓ edges, with μ changed by at most $\ell\Delta$ (`depPath_stepInvariant_with_sink`). Since the sink reads nothing in V_0, every arc $q \to p$ inside V_0 has $|\psi(p) - \varphi(q)| \le c_0 + D\Delta + c_1$: it is near, and Lemma 3.7(b) applies. □

Lemmas 3.7 to 3.9 are statements about traces only. They are the tools behind the exclusions of the schedule families in Appendix A; Theorem 5.1 is a far-reaching use of Lemma 3.8.

4 Two long-range heads

The first universal consequence of the certificate method concerns how far the heads must travel. It has polynomial constants and holds for every trace that carries a fixed exponentially small fraction of the inputs.

For reals $c, \beta \ge 1$ and an integer $k \ge 1$ put

$$N = \lfloor 160\beta c^2 \rfloor, \qquad \varepsilon_* = \frac{1}{8(N+1)}, \qquad \varepsilon_{\mathrm{sm}} = \frac{\varepsilon_*}{4\beta(N+2)\max(1,2k)},$$

$$d_{\mathrm{H}} = \max\Big(\lceil 80\beta c \rceil,\ \Big\lceil \frac{20(N+2)}{\varepsilon_*} \Big\rceil,\ \Big\lceil \frac{16\beta(N+2)k}{\varepsilon_*} \Big\rceil,\ 10\Big)$$

(`slabWallCap`, `slabAutonomyEps`, `slabSmallExcursion`, `slabDimThreshold`). For $c = 1$, $\beta = 2$ and $k = 2$: $\varepsilon_* \approx 3.9 \cdot 10^{-4}$, $\varepsilon_{\mathrm{sm}} \approx 3.8 \cdot 10^{-8}$ and $d_{\mathrm{H}} \approx 5.3 \cdot 10^{7}$.

Theorem 4.1 (Theorem H, two long-range heads)**.** Lean: `oneFreeHead_smallExcursion_excludes`, `twoMacroscopicHeads_of_linearTimeTransposer`, `twoMacroscopicHeads_of_oblivious_linearTimeTransposer`, `offCarrier_vertexExpansion_of_linearTimeTransposer`, `offCarrier_vertexExpansion_of_obliviousLinTime` [`OneFreeHeadSmallExcursion.lean`]

Let $c, \beta \ge 1$ and $k \ge 1$. Let M have at most k work tapes and value width at most β, let $d \ge d_{\mathrm{H}}$, and suppose M transposes at size d within time c. Let τ be an η-heavy trace of M with halting time T; for an oblivious M, take its common trace and $\eta = 0$.

(a) Slab autonomy. *If $\eta < \varepsilon_*$, there is no set of at most one work head such that every work head outside it has all its revisits of excursion at most $\lfloor \varepsilon_{\mathrm{sm}} L \rfloor$ on τ. The input head and the excepted head are unrestricted.*

(b) Two heads. *If $\eta < \varepsilon_*$, two distinct work heads each make, on τ, a revisit of excursion greater than $\varepsilon_{\mathrm{sm}} L$.*

(c) Off-carrier expansion. *For every set S of at most one work head and every $H \le \lfloor \varepsilon_{\mathrm{sm}} L \rfloor$, there is a set V_* of input positions with $3\varepsilon_* L \le |V_*| \le L/2$ such that every set B of steps separating*

$u_0(V_*)$ from the earlier ends of revisits of excursion greater than H on the work heads outside S satisfies

$$(\varepsilon_* - \eta)L \;\le\; \beta|B| + |B \cap u_0(V_*)|.$$

Part (b) is part (a) applied to the set of heads that make long revisits. In (c) the set V_* depends on the trace and on S; unlike Theorem X, (c) does not hold for every large output set, and its targets are all long revisits off one head, not only crossing ones. It has no hypothesis on η, but is vacuous for $\eta \ge \varepsilon_*$.

Outline of the formal proof. Frozen-trace chunks. If every dependency path that starts in a set C of steps and avoids Γ after its first vertex stays in C, then the input positions whose output copies are appended in C, minus those read in C, form an independent set U_*, and counting the fibre as in Lemma 3.6 gives $|U_*| - \beta|\Gamma| \le \eta L$ (`frozenTrace_chunk_deficit_le`). *Slabs.* Cut the tape of the unrestricted head into slabs of width about $L/(80\beta c)$; some offset makes it cross slab walls at most N times (`exists_offset_slabTransitions_le`). *A chunk.* Some time interval inside one slab epoch, of at most $L/2$ steps, appends at least $3\varepsilon_* L$ output bits (`exists_slab_chunk_dichotomy`). *Trapping.* A genie made of the slab transitions, the last visits of the free head before it leaves the slab, and the earlier ends of the other heads' revisits that span the chunk boundaries traps every path from the chunk (`depPath_trapped_in_chunk`); the last part is small because those revisits are short, and a short revisit can span a given step on at most $2H+1$ cells (`card_smallCrossingLowers_le`). *Counting.* The chunk has at most $N+1$ runs, and every output row meets every input row in b bits, so the independent set keeps about half the outputs of the chunk (`outIn_card_le_two_indep_add_six_runs`). The resulting deficit $\varepsilon_* L$ exceeds ηL, a contradiction. For (c), keep in the genie only the short revisits, which are always few, and add the separating set B: paths that avoid both stay in the chunk, and the same count gives the inequality. □

Corollary 4.2 (Sublinear excursions, family heights, one tape)**.** Lean: `oneFreeHead_sublinearExcursion_excludes`, `oneFreeHead_sublinearExcursion_oblivious_excludes`, `oneFreeHead_sublinearFamHeight_excludes`, `atMostOneWorkTape_lowAdaptivity_excludes`, `atMostOneWorkTape_oblivious_excludes` [`OneFreeHeadSmallExcursion.lean`]

Let $c, \beta \ge 1$, $k \ge 1$ and $H(d) = o(L)$. For all large d, no machine with at most k work tapes and value width at most β that transposes at size d within time c has an η-heavy trace, $\eta < \varepsilon_*$, on which all but one work head make only revisits of excursion at most $H(d)$; the same holds with the excursion replaced by the height of every family of revisits that share a side and a turning point. In particular no such machine has at most one work tape.

Outline of the formal proof of Corollary 4.2. For d large, $H(d) \le \varepsilon_{\mathrm{sm}} L$, so Theorem 4.1(a) applies; the height of a family bounds the excursions of its revisits (`hopExcursionLEOn_of_famd_le`). With at most one work tape, take the excepted set to contain every work head. □

Superlinear lower bounds for one work tape are classical in the one-tape models [17, 16], and in the present model the fourth paper's analysis of macroscopic limits already excludes one work tape [7, Corollary 3.18(c)]; we include the case only because it falls out. The fourth paper also showed, by a written proof, that in every macroscopic limit of a linear-time transposer at least two work-head limits are non-monotone [7, Corollaries 3.18 and 3.24]. Theorem H(b) is a quantitative form of this at a single size, with explicit polynomial constants, on every η-heavy trace with $\eta < \varepsilon_*$. Theorem H(a) excludes one free head even when the other heads make excursions up to a small fixed fraction of L, which the limit statement does not exclude, and part (c) is new. Every oblivious

linear-time transposer therefore has at least two work tapes on which information is stored and retrieved across distances that are a fixed fraction of L.

5 Macroscopic crossings

This section proves the dichotomy for arbitrary linear-time traces (Theorem D) and its consequence for oblivious transposers (Theorem X). Both rest on one implication, proved for every trace without any hypothesis on what the machine computes: if a small set of steps cuts a large set of outputs off from the crossing steps CL_H, then the remaining cone relation contains a linear-size acyclic set. Section 3 has shown that such an acyclic set is incompatible with transposition.

5.1 Constants

Fix $k \geq 1$, reals $C \geq 1$ and $\alpha > 0$, and a depth $D \geq 0$. Put

$$m_0(k) = (2k+1)(k+1)^{2k}, \qquad K_{\mathrm{row}} = \lceil 24C/\alpha \rceil, \qquad B_{\mathrm{time}} = \lceil 8m_0C/\alpha \rceil,$$

$$\rho_0(k,C,\alpha) = \frac{1}{3\binom{m_0+2K_{\mathrm{row}}}{m_0}}, \qquad N_{\mathrm{thr}} = (2k{+}1){+}m_0(2k{+}1)(B_{\mathrm{time}}{+}1), \qquad E_0(k,D,C,\alpha) = (k{+}2)^D N_{\mathrm{thr}}(5k{+}1),$$

$$\theta_{\mathrm{E}}(k,D,C,\alpha) = \frac{\alpha}{12\,E_0(k,D,C,\alpha)}, \qquad d_{\mathrm{pb}}(k,C,\alpha) = \max\Big(K_{\mathrm{row}}B_{\mathrm{time}}, \Big\lceil \frac{12m_0(2B_{\mathrm{time}}+3)}{\alpha} \Big\rceil\Big).$$

Here $m_0(k)$ counts the *phase-exit chains* of the proof below; $m_0 = 12$, 405 and 28672 for $k = 1, 2, 3$. The constant ρ_0 does not depend on D or on the alphabet. For a machine M with k work tapes and value width w_{val}, and $\alpha > 0$, put

$$c_{\mathrm{E}}(M,C,\alpha) = \frac{\rho_0(k,C,\alpha)}{2w_{\mathrm{val}}}, \qquad D_{\mathrm{E}}(M,C,\alpha) = \mathrm{genieDepth}\big(k, w_{\mathrm{val}}, C, \alpha\rho_0(k,C,\alpha)/4\big),$$

$$\theta_{\mathrm{mac}}(M,C,\alpha) = \theta_{\mathrm{E}}\big(k, D_{\mathrm{E}}(M,C,\alpha/2), C, \alpha/2\big) = \frac{\alpha}{24\,E_0\big(k, D_{\mathrm{E}}(M,C,\alpha/2), C, \alpha/2\big)},$$

where genieDepth is the depth bound of the fourteenth paper's genie (Lemma 3.1). In Lean these are `conjectureE_m0`, `conjectureE_bridge_rho0` (through `timePullbackRhoMin`), `conjectureE_Nthresh`, `conjectureE_E0`, `conjectureE_theta`, `timePullback_gen_d0`, `conjectureE_cE`, `conjectureE_DE` and `conjectureE_thetaMacro`.

These constants are positive and astronomically small. For $k = 2$, $C = 1$: $1/\rho_0(2,1,1) \approx 10^{65.8}$ and $1/\rho_0(2,1,\frac{1}{2}) \approx 10^{105.4}$; for $k = 1$ and $k = 3$ at $C = \alpha = 1$ the values are $10^{12.6}$ and $10^{153.4}$. The threshold θ_{mac} is far smaller, because E_0 grows like $(k+2)^D$ and the genie depth D_{E} is itself at least $32w_{\mathrm{val}}C/(\alpha\rho_0)$: for $k = 2$ and $C = \alpha = 1$, $\theta_{\mathrm{mac}} \leq 4^{-10^{107}}$. Nothing in the problem suggests that the true values are this small (Remark 5.8).

5.2 The time-pullback bound

The first ingredient turns a structure in time into an acyclic set. Split $[0,T)$ into B_{time} consecutive *time blocks* $T_0, \ldots, T_{B_{\mathrm{time}}-1}$ of length $\lfloor T/B_{\mathrm{time}} \rfloor$, the last one extended to end at T.

Theorem 5.1 (Time pullback)**.** Lean:
`theorem_51_2_general_zbad_mais` [`MacroscopicThreshold.lean`]

Let τ be a trace with k work tapes that appends exactly $L = d^2b$ bits before step $T \leq C(L+1)$, with $d, b, m_0 \geq 1$, $C \geq 1$ and $\alpha > 0$, and let $d \geq d_{\mathrm{pb}}$ computed with this m_0. Let $V_0 \subseteq [0,L)$ with

$|V_0| \geq \alpha L$, let $Z \subseteq [0, L)$ with $|Z| \leq |V_0|/12$, let B be a set of steps and $D \geq 0$. Suppose there are at most m_0 indices ν, each with a nondecreasing sequence $\xi_\nu(0) \leq \cdots \leq \xi_\nu(B_{\mathrm{time}})$ and a permutation π_ν of the blocks, such that the following holds for every $q \in V_0 \setminus Z$: whenever a dependency path from $u_0(q)$ with at most D edges, whose vertices after the root avoid B, ends at a step of block T_β at which the input head reads a position of V_0, the root $u_0(q)$ lies in $[\xi_\nu(\pi_\nu(\beta)), \xi_\nu(\pi_\nu(\beta)+1))$ for some ν. Then the ball relation of radius D avoiding B contains an acyclic set $U \subseteq V_0$ with

$$|U| \;\geq\; \frac{|V_0|}{3\binom{m_0+2K_{\mathrm{row}}}{m_0}}.$$

Outline of the formal proof. For each q let $A(q)$ be the set of blocks β such that $u_0(q)$ lies in an interval assigned to β; since the intervals of one ν are disjoint, $|A(q)| \leq m_0$. For an input position p let $E(p)$ be the set of blocks in which p is read. The hypothesis says that if p is in the ball of q then $A(q)$ and $E(p)$ meet. A kinematic count (`kinematic_goodInputRows_card_ge`) shows that all but at most $3T/(K_{\mathrm{row}}W)$ input rows are touched by the input head in at most K_{row} runs, each shorter than a block, so $|E(p)| \leq 2K_{\mathrm{row}}$ off a quarter of V_0. Pairs with $A(q) \cap E(q) \neq \emptyset$, self-collisions, are bounded without any anti-concentration assumption, using the row and column structure of tpos and the monotonicity of append steps: for a fixed ν each threshold splits at most one output row (`card_selfCollisions_singleFamily_le`). Off the bad sets the relation is contained in the *block relation* $A(q) \cap E(p) \neq \emptyset$ with $|A| \leq m_0$, $|E| \leq 2K_{\mathrm{row}}$ and no self-loops, and a uniform block cut (`exists_isAcyclicIn_card_ge_of_uniformBlockCut`) extracts an acyclic set of the stated size. □

Theorem 5.1 is purely about the trace. The constant $1/\binom{m_0+2K_{\mathrm{row}}}{m_0}$ is the price of the uniform block cut over all possible block sets, and it is the source of the size of ρ_0.

5.3 The bridge

The second ingredient finds the time-pullback structure whenever CL_H can be cut off. We need a separation notion that ignores the roots themselves.

Definition 5.2 (Separation within depth D)**.** Lean: `IsBallPathSeparator` [`EnvelopeCrux.lean`]
Let V_0 be a set of input positions and S, B sets of steps. *B separates V_0 from S within depth D* if every dependency path from a root $u_0(q)$, $q \in V_0$, with at most D edges and with all its vertices after the root outside B, also has all its vertices after the root outside S (`IsBallPathSeparator`).

If B separates $u_0(V_0)$ from S in the sense of Definition 2.4, then B separates $V_0' = \{q \in V_0 : u_0(q) \notin B\}$ from S within every depth (`isBallPathSeparator_of_isMengerVertexCut`).

Theorem 5.3 (Bridge)**.** Lean: `conjectureE_bridge_exists_isAcyclicIn_crossingLowerGT2H`, `conjectureE_bridge_exists_isAcyclicIn_crossingLowerH` [`MacroscopicThreshold.lean`]
Let τ be a trace with k work tapes that appends exactly $L = d^2 b$ bits before step $T \leq C(L+1)$, with $C \geq 1$ and $\alpha > 0$. Let $D \geq 0$, let $1 \leq H \leq \theta_{\mathrm{E}}(k, D, C, \alpha)\, L$, and let $d \geq d_{\mathrm{pb}}(k, C, \alpha)$. If V_0 has at least αL elements and B separates V_0 from CL_H within depth D, then the ball relation of radius D avoiding B contains an acyclic set $U \subseteq V_0$ with $|U| \geq \rho_0(k, C, \alpha)|V_0|$.

Outline of the formal proof. Two cuts per tape. On every unit-speed walk there are two steps ϑ_+, ϑ_- such that every revisit of excursion greater than H that crosses no other such revisit straddles one of them: the upper ones straddle ϑ_+, the lower ones ϑ_- (`exists_two_cuts_covering_nonCrossingLargeHops`). *Dichotomy along paths.* Along a dependency path from V_0 that avoids B

after the root, every edge is a state edge, a revisit edge of excursion at most $2H$, or a revisit edge whose revisit crosses nothing large and therefore straddles one of the $2k$ cut steps; the separation hypothesis excludes large crossing revisits (`depPath_nonCrossing_dichotomy_of_isBallPathSeparator_crossingLowerH`). *Phases and chains.* The cut steps split time into at most $2k+1$ work phases; a path visits the phases along one of at most $m_0(k)$ phase-exit chains (`card_chainIdx_le`), and along a chain each passage between phases is a strictly antitone partial bijection between times (`crossingCut_nonCrossing_strictAnti`). Pulling the block boundaries back along each chain gives, for every chain, monotone root intervals of both parities, which is the structure required by Theorem 5.1 with $m_0 = m_0(k)$. *Guards.* Small revisits and the state edge are handled by a guard set of at most $(5k+1)H$ steps per threshold (`card_microHopStepGuard_le`); the roots that reach a guard within D edges form the exceptional set Z, of size at most $(k+2)^D N_{\mathrm{thr}}(5k+1)H = E_0 H \le \alpha L/12$ by the choice of H (`card_ancestors_le`). Theorem 5.1 applies. □

The second Lean declaration is the same statement for banded crossing sets with an explicit clause on overshooting revisits; for the unbanded set CL_H the clause holds for every trace (`clauseE1H_uncapped`). The form with $H = W$ is `conjectureE_bridge_exists_isAcyclicIn_crossingLower`.

5.4 Theorem D: a dichotomy for every linear-time trace

Theorem 5.4 (Theorem D)**.** Lean: `linearTimeTrace_crossingLowerGT2H_menger_dichotomy`, `linearTimeTrace_crossingLowerGT2H_ballSep_dichotomy` [`MacroscopicThreshold.lean`]

Let τ be any trace with k work tapes that appends exactly $L = d^2 b$ bits before step $T \le C(L+1)$, with $C \ge 1$ and $\alpha > 0$. Let $D \ge 0$, $1 \le H \le \theta_{\mathrm{E}}(k, D, C, \alpha/2)\, L$, $d \ge d_{\mathrm{pb}}(k, C, \alpha/2)$, and $c \le \frac{1}{2}$. For every set V_0 of at least αL input positions, exactly one of the following holds:

(i) some set B of fewer than $c|V_0|$ steps separates $u_0(V_0)$ from CL_H, and the ball relation of radius D avoiding B contains an acyclic set $U \subseteq V_0$ with $|U| \ge \frac{1}{2}\rho_0(k, C, \alpha/2)|V_0|$;

(ii) every set separating $u_0(V_0)$ from CL_H has at least $c|V_0|$ elements.

Outline of the formal proof. If (ii) fails, take a separating set B with $|B| < c|V_0| \le |V_0|/2$. Since u_0 is injective, at least half of the roots of V_0 lie outside B; the set V_0' of these has at least $(\alpha/2)L$ elements, and B separates it from CL_H within depth D. Theorem 5.3 at density $\alpha/2$ gives U. □

Theorem D mentions neither a machine nor transposition. Its content is the implication from (i)'s first half to its second: a cheap cut from the crossing steps leaves a large acyclic set. Read the other way, it says that a trace whose cone relations have no large acyclic set after cheap cuts must reach CL_H with linear expansion from every large output set. The transposition hypothesis enters only through the certificate method.

5.5 Theorem X

Theorem 5.5 (Theorem X, macroscopic crossings)**.** Lean: `conjectureE_obliviousLinTime_macroscopicCrossingLowers_menger_lb`, `conjectureE_obliviousLinTime_crossingLowerGT2H_menger_lb` [`MacroscopicThreshold.lean`]

Let M be a machine with k work tapes and value width w_{val}, and let $C \ge 1$ and $\alpha > 0$. Put $\theta = \theta_{\mathrm{mac}}(M, C, \alpha)$ and

$$d_0 = \max\Big(1,\ d_{\mathrm{pb}}(k, C, \alpha/2),\ \Big\lceil \frac{24 E_0(k, D_{\mathrm{E}}(M, C, \alpha/2), C, \alpha/2)}{\alpha} \Big\rceil,\ \Big\lceil \frac{2}{\theta} \Big\rceil \Big).$$

Let $d \geq d_0$, suppose ObliviousLinTime(M, d, C, T)*, and let $H^* = \lfloor \theta L \rfloor$. Then $W \leq H^*$ and $\theta L \leq 2H^*$, and for every set V_0 of at least αL input positions, every set of steps separating $u_0(V_0)$ from CL_{H^*} in the dependency DAG of the common trace has at least*

$$\tfrac{1}{2}\, c_{\mathrm{E}}(M, C, \alpha/2)\, |V_0| \;=\; \frac{\rho_0(k, C, \alpha/2)}{4 w_{\mathrm{val}}}\, |V_0|$$

elements. The same bound holds for CL_H at every $1 \leq H \leq \theta L$ and every $d \geq \max(1, d_{\mathrm{pb}}(k, C, \alpha/2))$.

Every revisit counted in CL_{H^*} has excursion greater than $2H^* \geq \theta L$: these are crossings of macroscopic revisits. When $\alpha \leq 1$, taking V_0 to be a set of $\lceil \alpha L \rceil$ input positions and the separating set to be CL_{H^*} itself gives the following.

Corollary 5.6 (Linearly many macroscopic crossings)**.** `Lean`: `conjectureE_obliviousLinTime_macroscopicCrossingLowers_menger_lb` [`MacroscopicThreshold.lean`]

Under the hypotheses of Theorem 5.5, if moreover $\alpha \leq 1$, the common trace has at least $\frac{\alpha \rho_0(k,C,\alpha/2)}{4 w_{\mathrm{val}}} L$ steps that are earlier ends of crossing revisits of excursion greater than θL.

Outline of the formal proof of Theorem 5.5 and Corollary 5.6. Thresholds. From $d \geq d_0$ one gets $W \leq \theta L$, so $1 \leq W \leq H^*$, and $\theta L \geq 2$, so $\theta L < H^* + 1 \leq 2H^*$. *Reduction.* Let B separate $u_0(V_0)$ from CL_{H^*}. If $|B| \geq |V_0|/2$ we are done, since $\frac{1}{2} c_{\mathrm{E}} \leq \frac{1}{24}$. Otherwise the set V_0' of $q \in V_0$ with $u_0(q) \notin B$ has at least $|V_0|/2 \geq (\alpha/2)L$ elements, and B separates it from CL_{H^*} within the depth $D_{\mathrm{E}}(M, C, \alpha/2)$. *Bridge.* By the definition of θ_{mac}, $H^* \leq \theta_{\mathrm{E}}(k, D_{\mathrm{E}}, C, \alpha/2)L$, and Theorem 5.3 at density $\alpha/2$ gives an acyclic $U \subseteq V_0'$ in the ball relation with $|U| \geq \rho_0 |V_0'|$. *Coding.* Lemma 3.5 with the genie of cost $(\alpha\rho_0/8)L$ and depth D_{E} gives $w_{\mathrm{val}}|B| \geq |U| - (\alpha\rho_0/8)L \geq (\rho_0/2)|V_0'|$, hence $|B| \geq c_{\mathrm{E}}|V_0'| \geq \frac{1}{2} c_{\mathrm{E}} |V_0|$. □

Remark 5.7 (The earlier, row-scale form)**.** With $H = W$ the threshold condition is automatic for large d, and the conclusion holds with CL_W, the crossing steps of revisits longer than two rows (`conjectureE_obliviousLinTime_crossingLowerGT2W_menger_lb`); the same bound holds with the larger set of both ends of revisits that cross or share a cell and a turning point with another revisit (`conjectureE_obliviousLinTime_mixGT2W_menger_lb`). Theorem X is stronger: raising the threshold from $2W$ to θL shrinks the target set, so rearrangements below a fixed fraction of L can never substitute for macroscopic crossings.

Remark 5.8 (Numerical)**.** We computed exact minimum separating sets, by maximum flow, from large output sets to the earlier ends of crossing revisits of excursion greater than E, for $E \in \{2W, L/64, L/16, L/4, L/2, L\}$, on several explicit oblivious schedules at $d \in \{32, 64, 128\}$ and $b = 1$. On the five linear-time cascade schedules of Section 10.3, and on two superlinear transposers, the ratio of the minimum separating set to $|V_0|$ was exactly 1 for every $E \leq L/2$. At $E = L$ three of the cascades make no such crossings at all, so the empty set separates, while heuristic bounds on their largest acyclic sets fell towards 0 as the radius grew. So, numerically, the threshold in Theorem D appears not to be a universal constant such as $\frac{1}{2}$. These measurements are far below the sizes at which the theorems apply and are no evidence for them; they suggest only that the true expansion constant may be of order 1 at thresholds of order L/C.

6 Skeletons: where the crossings must be

Theorem X says how much information must reach macroscopic crossings. This section says where they can be. At a scale that is a fixed fraction of L, a linear-time head can turn only boundedly

often, so its coarse motion is one of finitely many words; crossings force two turns with overlapping legs; and the expansion of Theorem X can always be found on one leg of one tape. A leg is monotone only at the coarse scale: inside it the head may move back and forth over distances up to the scale.

6.1 Crossings are zigzags

Lemma 6.1 (Crossing revisits are zigzags)**.** Lean: `crossing_isHop_opposite_polarity`, `exists_crossingZigzag_of_crossingHopsGT2H`, `card_tapeCrossingLowers_ge_of_macroOverlap` [`MacroscopicSkeleton.lean`]

Let y be a unit-speed walk with horizon T and $H \geq 1$.

(a) Two crossing revisits (v_1, u_1), (v_2, u_2) with $v_1 < v_2 < u_1 < u_2$ visit different cells and lie on opposite sides of them: one stays strictly above its cell and the other strictly below.

(b) If two crossing revisits both have excursion greater than $2H$, there are a sign $\varepsilon \in \{\pm 1\}$ and steps $s_0 < s_1 < s_2 < s_3 < T$ such that $h = \varepsilon y$ satisfies

$$h(s_1)-h(s_0) \geq 2H+1, \quad h(s_1)-h(s_2) \geq 2H+1, \quad h(s_3)-h(s_2) \geq 2H+1, \quad \min(h(s_1), h(s_3))-\max(h(s_0), h(s$$

(c) Conversely, for every ε and $s_0 < s_1 < s_2 < s_3 < T$, the walk has at least $2N$ steps that are earlier ends of crossing revisits of excursion greater than $2H$, where $N = \max\bigl(0, \min(h(s_1), h(s_3))-\max(h(s_0), h(s_2)) - 2H\bigr)$, whatever the walk does between these steps.

Outline of the formal proof. (a) A walk with unit steps cannot pass a cell without visiting it (`isWalk_ivt`); a revisit therefore stays on one side of its cell, and the side of the second revisit is forced by the fact that the first cell is visited inside it but its own cell is not. (b) Take the extreme points of the two excursions. (c) Every level strictly inside the overlap is crossed upward between s_0 and s_1, downward between s_1 and s_2 and upward again between s_2 and s_3; the last visits before the turns give, for each level, one upper and one lower revisit of excursion greater than $2H$, and these cross. □

6.2 Skeletons

Definition 6.2 (Skeleton at scale Λ)**.** Lean: `milestones`, `tapeSkeleton`, `macroTurns`, `traceSkeletonTuple` [`MacroscopicSkeleton.lean`]

Let y be a walk with $y(0) = 0$ and let $\Lambda \geq 1$. Scan the steps $t = 0, 1, \ldots, T-1$ keeping a grid point $g \in \mathbb{Z}$, initially 0. Whenever $|y(t) - g\Lambda| \geq \Lambda$, record a *milestone* (t, σ) with $\sigma = +$ if $y(t) > g\Lambda$ and $\sigma = -$ otherwise, and move g one step in the direction σ. The *skeleton* of y at scale Λ is the word $\sigma_1\sigma_2\cdots$ of recorded signs (`tapeSkeleton`). A *leg* is a maximal run of equal signs, and the number of *turns* is the number of sign changes (`macroTurns`). The skeleton of a trace is the k-tuple of the skeletons of its work tapes.

Lemma 6.3 (Confinement and finiteness)**.** Lean: `epochGrid_confines_and_exact`, `tapeSkeleton_length_mul_le`, `card_boundedSkeletonTuple`, `traceSkeletonTuple_mem_boundedSkeletonTuple` [`MacroscopicSkeleton.lean`]

(a) At every step $|y(t) - g(t)\Lambda| < \Lambda$, where $g(t)$ is the current grid point, and $y(t) = g(t)\Lambda$ exactly at each milestone. A skeleton of length m has $m\Lambda \leq T$.

(b) Let τ be any trace with k work tapes and $T \le C(L+1)$, with $d, b \ge 1$, $C \ge 1$, $\theta > 0$ and $d \ge \lceil 8/\theta \rceil$. At scale $\Lambda = \lfloor \theta L/4 \rfloor$ we have $\Lambda \ge 1$, and every work tape's skeleton has at most $M_{\max} = \lfloor 16C/\theta \rfloor$ letters. Hence the skeleton of τ belongs to the set of k-tuples of words of length at most $M_{\max}$ over $\{+,-\}$, which has exactly $(2^{M_{\max}+1}-1)^k$ elements and does not depend on d.

Outline of the formal proof. (a) is an induction over the scan. (b) Each milestone after the first needs the head to move Λ cells, so $m\Lambda \le T \le 2CL$, while $\Lambda \ge \theta L/8$ once $\theta L \ge 8$; hence $m \le 16C/\theta$. The count of words of length at most $M_{\max}$ is $\sum_{i \le M_{\max}} 2^i$. □

Definition 6.4 (Crossing-capable skeletons)**.** Lean: `IsScaleFreeCrossingCapable`, `CrossingCapableSkeletonTuple` [`MacroscopicSkeleton.lean`]

Let s be a skeleton and let $\gamma(i)$ be its grid coordinate after i letters. The skeleton is *crossing-capable* if for some sign ε and indices $i_0 < i_1 < i_2 < i_3 \le |s|$ the function $h = \varepsilon\gamma$ satisfies

$$h(i_1)-h(i_0) \ge 3, \quad h(i_1)-h(i_2) \ge 3, \quad h(i_3)-h(i_2) \ge 3, \quad \min(h(i_1),h(i_3))-\max(h(i_0),h(i_2)) \ge -1$$

(`IsScaleFreeCrossingCapable`): up three grid steps, down three, up three, with overlapping ranges. Such a skeleton has at least two turns.

Proposition 6.5 (Which skeletons carry crossings)**.** Lean: `isCrossingCapableSkeleton_of_tapeCrossingLowers_nonempty`, `isScaleFreeCrossingCapable_of_two_mul_le`, `two_le_macroTurns_of_isScaleFreeCrossingCapable`, `tapeCrossingLowers_nonempty_of_isDetourProofSkeleton` [`MacroscopicSkeleton.lean`]

Let y be a walk with $y(0) = 0$, $\Lambda \ge 1$ and $H \ge 2\Lambda$. If the tape of y has an earlier end of a crossing pair of revisits of excursion greater than $2H$, then its skeleton at scale Λ is crossing-capable; in particular it has at least nine letters and at least two turns. Conversely, if for some sign and indices $i_0 < i_1 < i_2 < i_3$ the skeleton's grid satisfies $(\min(h(i_1),h(i_3)) - \max(h(i_0),h(i_2)))\Lambda \ge 2H+1$, then the tape has such crossing steps.

Outline of the formal proof. The first part combines Lemma 6.1(b) with the confinement of Lemma 6.3(a): the four extreme steps lie within Λ of grid points, which turns the inequalities with margin $2H+1 \ge 4\Lambda+1$ into the grid inequalities above (`isCrossingCapableSkeleton_of_tapeCrossingLowers_nonempty`, which proves a finer statement at every scale). The converse is Lemma 6.1(c) applied at the milestone steps, where the walk is exactly on the grid. □

Between the two conditions there is a narrow band of skeletons that may or may not carry crossings, depending on the motion inside the legs; this band is non-empty, so no condition on the skeleton alone decides whether crossings occur.

Lemma 6.6 (No macroscopic revisit within a leg)**.** Lean: `legOfEpoch_lt_of_isHop_gt2H` [`MacroscopicSkeleton.lean`]

Let y be a walk with $y(0) = 0$, $\Lambda \ge 1$ and $H \ge 2\Lambda$. Every revisit of excursion greater than $2H$ starts in an earlier leg of the skeleton than the one in which it ends.

Outline of the formal proof. Within one leg the grid point moves monotonically. A revisit returns to its cell, so if it started and ended in the same leg it would be confined to at most two adjacent epochs, with excursion less than $3\Lambda \le 2H$. □

6.3 Theorem S

For an oblivious linear-time transposer M with time constant C and $\alpha > 0$, put $\theta = \theta_{\mathrm{mac}}(M, C, \alpha)$, $\Lambda^* = \lfloor \theta L/4 \rfloor$, $H^* = \lfloor \theta L \rfloor$ and $M_{\max} = \lfloor 16C/\theta \rfloor$, and for a work tape j and $r \geq 0$ let $\mathrm{CL}^{j,r}$ be the set of steps of CL_{H^*} that are earlier ends of revisits on tape j and lie in the r-th leg of tape j's skeleton at scale Λ^* (`legCrossingLowers`).

Theorem 6.7 (Theorem S, skeletons and a single leg)**.** Lean: `obliviousLinTime_reduces_to_crossingCapableSkeleton_and_singleLeg`, `obliviousLinTime_infinite_pinned_skeleton_and_leg` [`MacroscopicSkeleton.lean`]

Let M be a machine with k work tapes, $C \geq 1$ and $0 < \alpha \leq 1$, and let $d \geq \max(d_0, \lceil 8/\theta \rceil)$ with d_0 as in Theorem 5.5. Suppose $\mathrm{ObliviousLinTime}(M, d, C, T)$.

(a) $1 \leq \Lambda^*$ and $2\Lambda^* \leq H^*$, and the skeleton of the common trace at scale Λ^* belongs to the explicit finite family of k-tuples of words of length at most $M_{\max}$ in which some tape's word is crossing-capable.

(b) For every set V_0 of at least αL input positions there are a work tape j, whose skeleton is crossing-capable with at least nine letters and two turns, and a leg $r \leq M_{\max}$ such that every set separating $u_0(V_0)$ from $\mathrm{CL}^{j,r}$ has at least

$$\frac{c_{\mathrm{E}}(M, C, \alpha/2)}{2k(M_{\max} + 1)} |V_0|$$

elements.

(c) If $\mathrm{ObliviousLinTime}(M, d, C, T_d)$ *holds for all d in an infinite set S, then there are one tuple of words, one tape and one leg such that, for infinitely many $d \in S$, the skeleton is that tuple and (b) holds with that tape and leg for V_0 the set of all input positions (with $\alpha = 1$).*

Outline of the formal proof. (a) is Lemma 6.3(b) together with Proposition 6.5, since CL_{H^*} is nonempty by Corollary 5.6. (b) By Lemma 6.6 each step of CL_{H^*} lies in a slot (j, r) with $r \leq M_{\max}$, and there are at most $k(M_{\max} + 1)$ slots. If every slot had a separating set smaller than the stated bound, their union would separate $u_0(V_0)$ from all of CL_{H^*} (`isMengerVertexCut_biUnion`) with fewer than $\frac{1}{2} c_{\mathrm{E}} |V_0|$ elements, contradicting Theorem 5.5. (c) The skeleton, tape and leg take finitely many values, independent of d, so one of them recurs infinitely often (`exists_infinite_fiber_of_mapsTo_finset`). □

The tape and the leg in (b) may depend on V_0 and on d; (c) removes the dependence on d along a subsequence. The finite family in (a) is astronomically large, because $M_{\max} = \lfloor 16C/\theta \rfloor$ inherits the size of θ; what matters is that it does not grow with d. The Lean proves only that the crossing-capable tuples form a proper subfamily of all tuples (the idle tuple is excluded, `card_crossingCapableSkeletonTuple_lt`), not how much smaller it is.

Remark 6.8 (Numerical: what the skeleton does not see)**.** Inside a leg the head may move back and forth over distances up to Λ^*, which is a fixed fraction of L. The skeleton does not see this motion, and the explicit cascades of Section 10.3 hide all but boundedly many of their refinement stages inside it. Our measurements on those schedules, and on a one-pass shear that Theorem 9.4 excludes, found skeletons independent of d with a single leg carrying linear expansion ($d \leq 128$). So, as far as these measurements go, the skeleton condition does not separate these schedules from a transposer.

7 Rows to columns

Theorems H, X and S use transposition only through the certificate method: a transposer cannot have a large acyclic set after a cheap cut. This section uses the transposition demand directly. Information from a set of input rows must reach the output rows into which transposition sends it, and every cut between the two must carry it.

For sets $A_0, A_1 \subseteq \{0, \dots, d-1\}$ of rows let $\mathsf{In}(A_0)$ be the set of steps at which the input head reads a position in an input row of A_0 (`inSlabReadSteps`), and let $Q^*(A_0, A_1)$ be the set of input positions whose input row lies in A_0 and whose output row lies in A_1 (`bichromaticRect`): the bits that rows A_0 send into output rows A_1. It has exactly $|A_0||A_1|b$ elements (`card_bichromaticRect`).

Lemma 7.1 (Lemma R, rows to columns)**.** Lean: `isAcyclicIn_coneSystem_uncutRootedBichromaticRect`, `frozenTrace_bichromatic_tposOf_cut_ge`, `oblivious_bichromatic_tposOf_cut_ge` [`TranspositionCut.lean`]

Let M transpose at size $d \geq 1$ within time c, and let τ be an η-heavy trace of M; for an oblivious M take its common trace and $\eta = 0$. For all sets A_0, A_1 of rows and every set Γ of steps separating $u_0(Q^*(A_0, A_1))$ from $\mathsf{In}(A_0)$,

$$w_{\mathrm{val}}|\Gamma| + |\Gamma \cap u_0(Q^*(A_0, A_1))| \;\geq\; |A_0||A_1|b - \eta L, \qquad \textit{hence} \qquad |\Gamma| \;\geq\; \frac{|A_0||A_1|b - \eta L}{w_{\mathrm{val}} + 1}.$$

Outline of the formal proof. Let Q' be the set of $q \in Q^*(A_0, A_1)$ whose root is not in Γ; since u_0 is injective, $|Q'| \geq |A_0||A_1|b - |\Gamma \cap u_0(Q^*)|$. For $q \in Q'$, a dependency path from $u_0(q)$ that avoids Γ after its root cannot end at a read of rows A_0: the root is not in Γ, so the path would avoid Γ altogether. Hence the cone of every $q \in Q'$ avoiding Γ contains no position of rows A_0, while Q' lies in rows A_0. No element of Q' sees another, so Q' is acyclic, and Lemma 3.6 gives $|Q'| - w_{\mathrm{val}}|\Gamma| \leq \eta L$. □

With $A_0 = A_1$ the set of all rows, Lemma R says that every cut between the input and the output carries the whole input: $|\Gamma| \geq (1 - \eta)L/(w_{\mathrm{val}} + 1)$. With smaller row sets it says the same for every block of the transposition demand separately. The filtering of roots matters: $\Gamma = u_0(Q^*)$ is a separating set, and the cone relation avoiding it need not be acyclic on all of Q^*, because roots in Γ may still see their own inputs.

The lemma becomes a statement about a single set of steps through a colouring.

Theorem 7.2 (Colouring criterion)**.** Lean: `isMengerVertexCut_of_bichromaticColoring`, `frozenTrace_bichromaticColoring_wave_cut_ge` [`TranspositionCut.lean`]

Let M, d, c, η and τ be as in Lemma 7.1, with halting time T, and let A_0, A_1 be sets of rows. Let Y be a set of steps and χ a colouring of the steps by $\{0, 1\}$ such that

(i) $\chi = 0$ on $\mathsf{In}(A_0) \setminus Y$ and $\chi = 1$ on $u_0(Q^*(A_0, A_1)) \setminus Y$;

(ii) every revisit (u, t) with $u, t \notin Y$ has $\chi(t) \leq \chi(u)$.

Let $\partial\chi$ be the set of steps $t \in [1, T)$ with $t, t-1 \notin Y$, $\chi(t) = 1$ and $\chi(t-1) = 0$. Then every set B of steps separating the roots of all input positions from Y satisfies

$$|B| \;\geq\; \frac{|A_0||A_1|b - \eta L}{w_{\mathrm{val}} + 1} - |\partial\chi|.$$

Outline of the formal proof. $B \cup \partial\chi$ separates $u_0(Q^*)$ from $\mathsf{In}(A_0)$. Take a dependency path from a root to a read of rows A_0 that avoids $B \cup \partial\chi$. It avoids Y, since its part up to a vertex of Y would be a path from a root to Y avoiding B. By (i) the path starts with colour 1 and ends with colour 0, so some edge drops from colour 1 to colour 0 as the path moves to an earlier step. A state edge $t \to t-1$ with $\chi(t) = 1$ and $\chi(t-1) = 0$ would put t in $\partial\chi$, and a revisit edge $t \to u$ with $\chi(t) = 1$ and $\chi(u) = 0$ is excluded by (ii), a contradiction. Lemma 7.1 bounds $|B \cup \partial\chi|$. □

Taking $B = Y$ shows $|Y| \geq (|A_0||A_1|b - \eta L)/(w_{\mathrm{val}} + 1) - |\partial\chi|$: a set of steps across which the colour can only jump at few places must itself be large. Several such sets that are pairwise disjoint give a time bound. For $d \geq 16$ let $K = \lfloor \log_2 d \rfloor$ and $d_1 = 2^K$, and for $0 \leq s \leq K-2$ let A_s^0 and A_s^1 be the rows $i < d_1$ whose bit $K-1-s$ is 0, respectively 1.

Corollary 7.3 (Scale-separated traces)**.** Not formalized; a complete proof is given below.

Let M, d, c, η and τ be as in Lemma 7.1, with $d \geq 16$, $\eta \leq \frac{1}{32}$, $K_{\mathrm{flip}} \geq 0$ and $d \geq 64(w_{\mathrm{val}} + 1)K_{\mathrm{flip}}$. Suppose that for each s in a set $S \subseteq \{0, \ldots, K-2\}$ there are a set Y_s of steps and a colouring χ_s satisfying (i) and (ii) of Theorem 7.2 for (A_s^0, A_s^1) with $|\partial\chi_s| \leq K_{\mathrm{flip}} d$, and that the sets Y_s are pairwise disjoint. Then $T \geq |S| \cdot L/(64(w_{\mathrm{val}} + 1))$.

Proof (not formalized). Since $d_1 > d/2$, $|A_s^0||A_s^1|b = (d_1/2)^2 b > L/16$. Since $d \geq 64(w_{\mathrm{val}} + 1)K_{\mathrm{flip}}$ and $b \geq 1$, $K_{\mathrm{flip}} d \leq d^2/(64(w_{\mathrm{val}} + 1)) \leq L/(64(w_{\mathrm{val}} + 1))$. Theorem 7.2 with $B = Y_s$ gives

$$|Y_s| \;\geq\; \frac{(1/16 - \eta)L}{w_{\mathrm{val}} + 1} - K_{\mathrm{flip}} d \;\geq\; \frac{L}{32(w_{\mathrm{val}} + 1)} - \frac{L}{64(w_{\mathrm{val}} + 1)} = \frac{L}{64(w_{\mathrm{val}} + 1)}.$$

The Y_s are pairwise disjoint subsets of $[0, T)$, so $T \geq \sum_{s \in S} |Y_s|$. □

With S all $K-1$ scales the bound is $T \geq (\lfloor \log_2 d \rfloor - 1)L/(64(w_{\mathrm{val}} + 1))$, so no linear-time transposer has a trace that is scale-separated across all scales. The corollary should be read with Proposition 10.2 in mind: one set of steps, together with the colouring that changes once, at the step after the last input read, satisfies (i) and (ii) for every pair of row sets at once, so the content of the hypothesis is the disjointness of the sets Y_s, which says roughly that live storage is rewritten once per scale. The corollary is therefore a time bound for schedules that refine the row index one bit at a time in separate passes, whatever they compute; it does not show that a linear-time transposer must do so.

8 Adaptive machines

An adaptive machine has many traces and no common dependency DAG. Two routes lead from the oblivious results to adaptive machines. The first uses a single heavy trace of the machine itself; the second passes through the bit simulator of the fifteenth paper, whose traces are few.

Theorem 8.1 (Theorem E, heavy traces and simulators)**.** Lean:
`frozenTrace_macroscopicCrossingLowers_menger_lb`,
`frozenTrace_reduces_to_crossingCapableSkeleton_and_singleLeg`,
`adaptive_simulator_heavyTrace_macroscopicExpansion`,
`simulator_heavyTrace_reduces_to_crossingCapableSkeleton_and_singleLeg`
[`SimulatorBridge.lean`]

(a) Heavy traces. *Let M be any machine with k work tapes, $1 \leq C$, $0 < c \leq C$ and $\alpha > 0$, and let $d \geq d_0$ with d_0 as in Theorem 5.5. Suppose M transposes at size d within time c,*

and let τ be an η-heavy trace of M with $\eta \le (\alpha/8)\rho_0(k, C, \alpha/2)$. Then the conclusion of Theorem 5.5 holds in the dependency DAG of τ, with the same constants; if moreover $\alpha \le 1$ and $d \ge \lceil 8/\theta_{\mathrm{mac}}(M, C, \alpha)\rceil$, so do the conclusions (a) and (b) of Theorem 6.7.

(b) Simulators. *Let M be any machine with k work tapes, $c > 0$ and $0 < \alpha \le 1$, and put $\eta^* = (\alpha/32)\rho_0(k, 11c+1, \alpha/2)$. There are $s \ge 1$ and d_{sim} such that the simulator N_s of [14, Theorem S], which has the same work tapes and work alphabet as M and computes the same partial function, satisfies the following at every $d \ge d_{\mathrm{sim}}$ at which M transposes within time c: N_s transposes within time $11c+1$; it has at most $2^{\eta^* L}$ distinct traces on the inputs of length L; its $2\eta^*$-heavy traces are followed together by more than $(1 - 2^{-\eta^* L})2^L$ inputs; and on each of its $2\eta^*$-heavy traces the conclusions of Theorem 5.5 and of Theorem 6.7(a),(b) hold in the dependency DAG of that trace, with the constants of N_s and the time constant $11c+1$.*

Outline of the formal proof. (a) The proof of Theorem 5.5 goes through on the fixed trace τ: Theorem 5.3 is a statement about traces, and in the coding step Lemma 3.6 replaces the index code of Lemma 3.1(b). With the genie of cost $(\alpha\rho_0/8)L$ and the heaviness $\eta \le (\alpha/8)\rho_0$, the two losses together are at most $(\rho_0/4)|V_0|$. (b) The constant η^* depends only on k, c and α. The fifteenth paper's trace transfer at η^* (`TapeStorageRates.aligned_low_trace_dimensions`) gives s and the simulator, with at most $2^{\eta^* L}$ traces; the light traces are followed by fewer than $2^{\eta^* L}2^{(1-2\eta^*)L}$ inputs. Since N_s has k work tapes and runs in time $11c+1$, $2\eta^* \le (\alpha/8)\rho_0(k, 11c+1, \alpha/2)$, and part (a) applies to N_s on each heavy trace. The quantifiers are ordered η^*, then s and N_s, then the constants of N_s and d_{sim}. □

Part (b) is the only route by which the macroscopic condition reaches arbitrary linear-time transposers so far, and it reaches them through a different machine's dependency DAG. The simulator rewrites windows of $3s$ cells in each round, and a separating set in its DAG need not correspond to one in the DAG of M (Problem 11.2). Its value width grows with s, so its constants are much smaller than those of M. Statements in this paper about the dependency DAG of a general machine therefore always assume a heavy trace or name the simulator.

Remark 8.2 (The heaviness hypothesis cannot simply be dropped)**.** Proof sketch only: given any machine M_0, add three states so that the new machine first scans the input for a 1; on the all-zero input it writes the all-zero output with its work heads standing still, and otherwise it returns its input head and runs M_0. The new machine transposes whenever M_0 does, within a constant more time, and its all-zero trace makes no revisits at all. Hence the version of Theorem 8.1(a) without a lower bound on the trace probability would exclude this machine, and so would imply ¬(T). It is not refuted: it is false if (T) holds.

Remark 8.3 (Wave counts on heavy traces)**.** Because every expanding set of steps is as large as its expansion, m pairwise disjoint sets of steps into which a set of αL outputs expands by a factor c need time at least $mc\alpha L$ (`disjointWave_mengerExpansion_counting_le`); for oblivious machines in time $C(L+1)$ there are at most $2C/(c\alpha)$ of them (`obliviousLinTime_expandingWaves_le`). A hypothesis asserting unboundedly many such sets on every heavy trace of every linear-time transposer is therefore equivalent to there being no linear-time transposer at infinitely many sizes, and gives no route to ¬(T). For oblivious machines, a hypothesis asserting unboundedly many such sets on every transposer, of any running time, would exclude oblivious linear-time transposers; the converse is not known.

9 Class exclusions

The results of this section exclude classes of machines defined by their head motion. They use the certificate method directly and reach machines that Theorems X and S do not: schedules with macroscopic crossings that are nevertheless excluded, and adaptive machines with no heavy trace.

9.1 Boundedly many macroscopic phases

A walk is *weakly monotone* on an interval of steps if it never moves left there, or never moves right (pauses allowed). It is *A-coarsely monotone* if it never moves more than A cells against a fixed direction: $x(t') \geq x(t) - A$ for all $t \leq t'$ in the interval, or $x(t') \leq x(t) + A$ for all such $t \leq t'$ (`CoarseMonoOnInterval`).

Definition 9.1 (Bounded macro-phases)**.** Lean: `WeakMonoOnInterval`, `BoundedMacroBank` [`MacroBankCrux.lean`]

Let M be a machine, $C \geq 1$, $M_{\mathrm{ph}} \geq 1$, B a set of work heads (the *bank*) and $S : \mathbb{N} \to \mathbb{N}$ with $S(d)/L \to 0$. M has *bounded macro-phases* with these parameters (`BoundedMacroBank`) if for all large d, whenever M is oblivious at length L with halting time $T \leq C(L+1)$, the interval $[0, T)$ splits into M_{ph} consecutive intervals on each of which the input head and every bank head are weakly monotone, and every other work head visits at most $S(d)$ cells before T.

The heads outside the bank are unrestricted inside their $o(L)$ workspaces: any number of reversals, any nesting, any coding.

Theorem 9.2 (Theorem B, bounded macro-phases)**.** Lean: `boundedMacroBank_has_crux`, `boundedMacroBank_excludes` [`MacroBankCrux.lean`]

If M has bounded macro-phases with bank B and M_{ph} phases, then M has crux rate $1/(4(|B|+1)^{M_{\mathrm{ph}}-1})$ at every depth, and for all large d there is no T with $\mathrm{ObliviousLinTime}(M, d, C, T)$.

Outline of the formal proof. Let $K_0 = (|B|+1)^{M_{\mathrm{ph}}-1}$ and cut workspace revisits longer than $\lceil\sqrt{S(T+1)}\rceil$; there are $o(L)$ of them. Inside a phase a bank revisit goes back by one step; between phases, bank revisits are strictly monotone partial injections of fixed sign. So after deleting $o_D(L)$ exceptional outputs, every short dependency path from an output has a *signature*, the decreasing sequence of phases it enters through bank heads, and at most K_0 signatures start in a given phase. Thresholds on input addresses pull back along signatures to thresholds in time. Split the input rows into $2K_0$ buckets: each unsplit output row forbids at most K_0 buckets, so some bucket is missed by half of the output rows, and the rectangle of that bucket and those output rows, minus the exceptions, is independent in the cone relation; it has $(1/(4K_0) - o(1))L$ elements (`macroBank_independent_exists`). □

Theorem 9.3 (Suffix phases)**.** Lean: `suffixMacroBank_has_crux`, `suffixMacroBank_excludes` [`SuffixMassCrux.lean`]

Theorem B holds, with crux rate $\delta_0/(4(|B|+1)^{M_{\mathrm{ph}}-1})$, for machines that satisfy the conditions of Definition 9.1 only on a final interval of steps that contains every append, provided at least $\delta_0 d$ input rows are not read before that interval. The motion before it is unrestricted.

Outline of the formal proof. Dependency paths decrease time, so a path from an output to a read of an untouched row stays inside the final interval (`getLastD_le_of_mem_depPath`); run the proof of Theorem B there, with the buckets taken among the untouched rows. □

For adaptive machines the phases and the workspaces may depend on the input, and the walks may jitter.

Theorem 9.4 (Theorem J, jittered phases, any adaptivity)**.** Lean: `native_machine_preserves_jitterBM`, `nativeGapBound_of_adaptiveJitterBM`, `adaptiveJitterBM_excludes_transposesAt`, `adaptiveJitterBM_not_linearTimeTransposer`, `adaptiveSuffixJitterBM_excludes_transposesAt`, `adaptiveSuffixJitterBM_not_linearTimeTransposer` [`JitterMacroBank.lean`,`AdaptiveSuffixJitterBM.lean`]

Let M be any machine, $c > 0$, and $A(d), S(d) = o(L)$. Suppose that for all large d every run of M on an input of length L that halts within $c(L+1)$ steps splits into M_{ph} intervals on each of which the input head and the bank heads are $A(d)$-coarsely monotone, while every other work head visits at most $S(d)$ cells; the intervals and cells may depend on the input. Then M does not transpose within time c at all large d; in particular M is not a linear-time transposer. The same holds if the conditions hold only on a final interval containing every append and at least $\delta_0 d$ rows are not entered by the input head before it.

Outline of the formal proof. The bit simulator of the fifteenth paper preserves the shape, with the jitter and the workspace enlarged by $O(s)$ (`native_machine_preserves_jitterBM`). On each frozen trace of the simulator, the proof of Theorem B goes through with jitter: the revisits that straddle a pulled-back cut on a bank head occupy at most $2A+1$ cells (`hop_upper_injOn`), and input reads are pulled back up to a band of width A. This gives every storage code on the simulator's frozen cone relations rate at most $1 - 1/(4K_0) + o(1)$ (`nativeGapBound_of_adaptiveJitterBM`), and Theorem 9.5 below, applied to the simulator, which has few traces, excludes M. □

At a fixed number of phases these classes probably also yield to the grouped pass bound of the first paper [9, Theorems 5.8 and 5.12] once the jitter and the workspaces are absorbed into the finite control; Remark 10.3 sketches why. What is new here is a kernel-checked proof with an explicit native coding deficit. The suffix forms are not reached that way: an unrestricted prefix may read $(1-\delta_0)d$ rows, and conditioning on the configuration at the start of the suffix leaves classes of inputs of co-entropy up to $(1-\delta_0)L$, for which the published robust pass bounds [9, 13] say nothing.

9.2 A native gap suffices

Theorem 9.5 (Theorem G, native gap criterion)**.** Lean: `nativeGapBound_half_iff`, `native_gap_tests_exclude_linear`, `native_gap_exists_exclude_linear` [`NativeGapCriterion.lean`]

Suppose that for every k and $C \geq 1$ there is $\rho(k, C) > 0$ such that, for every machine with k work tapes and every depth D, every binary storage code on the cone relation, after a genie of full depth D, of every run that halts within $C(L+1)$ steps on an input of length L (entries of $b(d)$ bits) has rate at most $1 - \rho(k, C) + o(1)$. Then (T) fails.

Outline of the formal proof. If (T) held, some machine M would transpose in time $c(L+1)$ at infinitely many sizes. Fix $C = 11c + 1$ and $\rho_* = \rho(k, C)$, and then choose the simulator N_s of [14, Theorem S] with at most $2^{(\rho_*/2)L}$ traces. Some trace of N_s carries $2^{(1-\rho_*/2)L}$ inputs; after the genie, its fibre splits into storage codes on its cone relation (`MergeDepth.blockMachine_isStorageCode_fibre`), one of which has rate at least $1 - \rho_*/2 - o(1)$, contradicting the hypothesis. At $\rho = \frac{1}{2}$ the hypothesis is the native half bound of the fourteenth paper [6, Corollary L], assumed uniformly over the machines with k work tapes and over all depths, as in the eighth paper [5, §6.2]. □

The fifteenth paper reaches the same conclusion from the half bound on chosen simulators [14]; Theorem G accepts any fixed deficit, uniformly over the machines with k work tapes. The hypothesis is at least as strong as a uniform crux; its value is that the deficit may be tiny, which is the form that Theorems B and J produce, but only for their classes.

9.3 Few scales and no free head

For a work head, a *family* is the set of its revisits on one side that share their *turning point*, the first step inside the revisit at which the head is farthest from the cell; its *height* is the distance from the turning point to the family's outermost cell, the cell of its longest revisit, as in the eighth paper [5]. The head has *owner load* at most ℓ if at every step, on each side, the revisits open at that step have at most ℓ distinct turning points. Bounded cell entry V gives owner load at most $2V - 1$.

Theorem 9.6 (Theorem Q, scale-gap pigeonhole)**.** `Lean: zeroFreeHead_qWindow_excludes, zeroFreeHead_entries_qWindow_excludes, zeroFreeHead_qWindow_oblivious_excludes` [`ZeroFreeHeadQWindow.lean`]

Let c, β be reals and k, ℓ, q, h integers. There are $\varepsilon > 0$ and d_0 such that for every $\eta < \varepsilon$ and $d \geq d_0$ no machine with at most k work tapes and value width at most β transposes at size d within time c with an η-heavy trace on which every work head has owner load at most ℓ and the binary logarithms of the heights of all families lie in at most q windows of width h. The input head is unrestricted. For oblivious machines the statement holds at every $d \geq d_0$ with $\eta = 0$.

Outline of the formal proof. Owner load ℓ colours the families with $N = 2k \max(\ell, 1)$ classes of non-overlapping spans, so their heights sum to less than NcL. Choose thresholds $Q_0 < Q_1 < \cdots < Q_q$, each so large compared with the last that a window can be *bad*, that is contain the height of a family longer than the current threshold, for at most one index. With q windows and $q + 1$ indices some index is bad for no window, so only boundedly many families are large, and the eleventh paper's low-adaptivity theorem for boundedly many large families [10] gives the contradiction. Its constant comes from a compactness argument and is not explicit [15]. □

Theorem Q is the case without a free work head of the eighth paper's bounded-load, boundedly-many-scales class, which is conditional on the crux there [5]; here it is unconditional, with any input motion and low adaptivity. With $\Theta(\log d)$ scales instead of boundedly many, the question is open already for two work tapes.

10 What a linear-time transposer needs

This section reads the results backwards, as a specification that any oblivious linear-time transposer must meet, and then describes how close known schedules come to meeting it.

10.1 The specification

Suppose M is an oblivious machine with k work tapes that transposes in time $C(L + 1)$ at infinitely many sizes. By the results of this paper and of the series, at all large such sizes:

(i) *Shape.* At scale $\theta L/4$ each work tape's walk is one of finitely many words, and some tape turns at least twice with overlapping legs, so it reuses a stretch of tape like a queue (Theorem 6.7). Along infinitely many sizes the shape is one fixed tuple of words.

(ii) *Crossings.* From every set of at least αL outputs, a linear number of vertex-disjoint dependency paths reach the earlier ends of crossing revisits of excursion more than θL (Theorem 5.5), and the crossing steps of one coarsely monotone leg of one tape already carry a linear fraction of them; the tape and the leg may depend on the output set and on the size, and are fixed along infinitely many sizes for the set of all outputs (Theorem 6.7).

(iii) *Two long-range heads.* At least two work heads revisit cells after moving a fixed fraction of L away, and for each work head some set containing a fixed fraction of the outputs reaches long revisits off that head with linear expansion (Theorem 4.1).

(iv) *Rows to columns.* For every two sets of rows A_0, A_1, every set of steps separating the reads of rows A_0 from the outputs that rows A_0 feed into output rows A_1 has at least $|A_0||A_1|b/(w_{\mathrm{val}}+1)$ elements (Lemma 7.1).

(v) *Codes.* After a genie of cost εL and bounded depth, the cone relations of the schedule carry an index code of at most εL bits, one of whose fibres is a storage code of rate at least $1-\varepsilon$ [6, Corollary O]; by (i)–(iii), the relations involved are those of macroscopic queue-like reuse.

Conditions (i)–(iv) are proved here without any coding hypothesis. Condition (v) is the coding question itself: a linear-time oblivious transposer is a code of rate close to 1 on the cone relations of a schedule meeting (i)–(iv). Machine-free relations carry storage codes above rate one half [11], and binary codes of rate approaching one exist on suitable separated relations [14].

10.2 Mechanisms that fail

Each of the following design patterns is excluded for oblivious linear-time transposition, whatever the machine computes on its tapes.

- *Stack-like reuse at macroscopic scale.* A schedule whose long revisits are nested, so that every tape reuses its storage like a stack at scale θL, has no macroscopic crossing steps, and is excluded by Theorem 5.5.
- *One long-range head.* One unrestricted work head together with heads whose revisits stay below a small fixed fraction of L is excluded by Theorem 4.1, whatever the input head does.
- *Boundedly many macroscopic phases.* Schedules in which the input head and some bank heads are weakly monotone on boundedly many phases while the other heads move inside sublinear workspaces are excluded by Theorem 9.2; with sublinear jitter on the bank heads, and for adaptive machines, by Theorem 9.4.
- *Cheap row-to-column cuts.* A schedule with a set of fewer than $|A_0||A_1|b/(w_{\mathrm{val}}+1)$ steps separating the reads of some rows A_0 from their outputs in output rows A_1 is excluded by Lemma 7.1; Theorem 7.2 gives a convenient way to exhibit such cuts.
- *The explicit families of Appendix A*, each excluded by a class theorem of Section 9 or a certificate of Section 3 (in proof sketches). Several of them make macroscopic crossings with linear expansion (numerically, at $d \le 128$), so Theorem 5.5 alone does not exclude them.

10.3 Schedules that meet every condition we could test

Five explicit oblivious schedules, which we call *two-sided cascades*, meet every necessary condition of this paper that we could test. Each loads the input onto two or three work tapes and then runs a cascade over the scales $M_r = L/2^r$, $1 \le r < \log_2 d$. A forward stage folds the data at each scale: one half by a transfer at varying relative speed, the other merged in with a cyclic shift of about $M_r/3$. A backward stage unfolds it by 1-to-2 doublings with an unshifted and a shifted half. A final sweep emits the L output bits. Every stage at scale r moves the heads over $O(M_r)$ cells a bounded number of times, so the running time is $O(\sum_r M_r) = O(L)$. The five differ as follows.

1. Two work tapes; before the cascade, two symmetric pairs of two-dimensional cyclic shifts of the entry grid are applied to a copy of the input.

2. Two work tapes; before the cascade, an anti-diagonal reflection and two butterfly shifts, by $2^{m-1}(d-1)b$ and $(d-1)b$ where $d = 2^m$, are applied to a copy.

3. Three work tapes, every cell entered at most ten times; the third tape holds disjoint zones for the folded data at all scales and moves through them monotonically in each stage.

4. Three work tapes; the forward and backward stages use disjoint zones for all scales, and no zone is overwritten.

5. Two work tapes; the folds are 2-to-1 compressions, and every scale has a comb pass in which one tape makes 2^r teeth while the other waits.

The machines that realise them in our experiments compute other functions; whether some machine with one of these oblivious traces transposes is Problem 11.4.

Remark 10.1 (Numerical)**.** At $d \le 128$ and $b = 1$ the measured running times of the five schedules lay between $18L$ and $90L$ and stabilised as d grew. On all five, the minimum set separating a large output set from the earlier ends of crossing revisits of excursion greater than E had exactly as many elements as the output set, for every tested $E \le L/2$. Their skeletons at scale about $L/6$ were independent of d, with a single leg carrying linear expansion, and their numbers of disjoint expanding sets of steps stayed bounded, as linear time requires. So, as far as these measurements go, the conditions (i)–(iv) of the specification do not distinguish these schedules from a transposer. A one-pass shear, excluded by Theorem 9.4, and a multi-scale stack schedule, excluded by a certificate (in a proof sketch, Appendix A), also make macroscopic crossings with linear expansion in these measurements, which suggests that (i)–(iv) are not sufficient for transposition.

10.4 Why cuts alone cannot finish

The certificates of this paper are cuts: they bound from below the number of steps that must separate inputs from outputs. A single cut can only certify storage, and the next proposition shows that for transposition one set of steps serves as the cut for every pair of row sets at once.

Proposition 10.2 (One cut for every pair of row sets)**.** Not formalized; a complete proof is given below. *Let a run with horizon T read its last input cell at step t_{in}, and put $t^* = t_{\text{in}} + 1$. Let W^* consist of*

- *every step $u < t^*$ that is the earlier end of a revisit (u, t) with $t \ge t^*$,*
- *every output step before t^*, and*
- *the step t^*, if $t^* < T$.*

Then for every two sets A_0, A_1 of rows, W^* meets every dependency path from the output steps of the entries that rows A_0 feed into output rows A_1 to the steps that read rows A_0. The steps of the first kind number at most the number of work cells that are visited both before t^* and at or after t^*.

Proof (not formalized). Take a dependency path from an output step t_0 to an input-reading step u. Every edge of the dependency DAG goes from a step to a strictly earlier step: the state edge $a \to a - 1$, or the revisit edge $a \to a'$ for a revisit (a', a). If $t_0 < t^*$, then t_0 is an output step before t^* and lies in W^*. Otherwise $t_0 \ge t^* > t_{\text{in}} \ge u$, so the path has an edge $a \to a'$ with $a \ge t^* > a'$. If

it is a state edge, then $a' = a - 1$, so $a = t^*$, and $t^* \in W^*$. If it is a revisit edge, then (a', a) is a revisit with $a' < t^* \le a$, so $a' \in W^*$. In each case the path meets W^*; this argument does not use A_0 or A_1. A step of the first kind is the earlier end of a revisit of some tape that spans t^*; on each tape, each cell is the cell of at most one such revisit, namely the one from its last visit before t^* to its next visit. □

For a transposer, Lemma 7.1 with $A_0 = A_1$ the set of all rows gives $|W^*| \ge (1 - \eta)L/(w_{\mathrm{val}} + 1)$. When few outputs precede the last input read, this says that the storage live across t^* is linear. In any case Proposition 10.2 shows that this one bound already accounts for every pair of row sets: however finely the rows are split, cut certificates for different splits may coincide, and a lower bound beyond linear time cannot come from counting disjoint cuts unless something other than the cut property forces the cuts apart. The colouring criterion of Section 7 verifies such a separation on specific schedules (Corollary 7.3), but it does not supply one in general.

Outside such restricted classes, the only coding-proof logarithmic factor known on tapes is of a different kind. For an oblivious transposer whose run has r monotone passes, Theorem P of the first paper gives $k\,db \le ((k+1)^r - 1)(2b + \sigma_M)$, where σ_M is the number of bits of a state and the scanned symbols, so a fixed machine needs $r \ge \log_{k+1} d - O(1)$ passes [9, Theorem 5.9]; for adaptive machines an extra term appears [9, Theorem 5.12]. Its grouped form [9, Theorem 5.8] reads $d^2 b \le H_{\mathrm{addr}}((\gamma+1)db + \lceil d/\gamma \rceil \sigma_M)$ for every $\gamma \ge 1$, where $H_{\mathrm{addr}} = ((k+1)^r - 1)/k$ bounds the number of addresses of the forests in its proof.

Remark 10.3 (Pass counting at one scale)**.** Proof sketch only: suppose a run splits into r phases in each of which every head stays within distance S of a monotone trend. Adapting the proof of the grouped bound so that the nodes of its forests carry the windows of $O(S)$ cells around the trends, at the cost of bookkeeping for the offsets within each phase, plays the role of a finite control of $\sigma \approx \sigma_M + O(kS \log(g+1))$ bits, and choosing $\gamma \approx \sqrt{\sigma/b}$ in the grouped inequality gives $(k+1)^r \gtrsim \frac{k}{2}\sqrt{L/\sigma}$. So boundedly many phases with $S = o(L)$ cannot transpose. At $S = \varepsilon L$, however, the bound only asks for $r \ge \frac{1}{2}\log_{k+1}(1/\varepsilon) - O(1)$ phases. A linear-time run may have up to about $4kC/\varepsilon$ legs at scale εL (Lemma 6.3), and schedules whose number of legs at scale εL grows like $\log(1/\varepsilon)$, such as the cascades of Section 10.3, meet the bound at every scale.

So pass counting, like cut counting, cannot exclude linear time one scale at a time. Problem 11.1 asks for a pass bound that counts phases across all scales at once.

11 Open problems

The results of this paper confine an oblivious linear-time transposer to macroscopic queue-like reuse with linear expansion, on one leg of a bounded skeleton, with two long-range heads. The following problems would take the next step; the first three are on the side of impossibility, the fourth on both sides.

Problem 11.1 (Pass bounds across scales)**.** *For a run and $1 \le j \le \log_2 L$ let r_j be the total number of legs of its skeletons at scale $\lfloor 2^{-j} L \rfloor$. A linear-time run has $r_j \le k(4C\,2^j + 1)$ (Lemma 6.3), and pass counting at a single scale requires only $r_j \gtrsim \frac{j}{2}\log_{k+1} 2$ (Remark 10.3). Is there, for every k and C, a function f with $f(d) \to \infty$ such that no oblivious machine with k work tapes whose runs satisfy $r_j \le Cj$ for all j transposes at size d in time less than $f(d)L$? More generally, find a coding-proof inequality that involves all the r_j at once and fails for every linear-time run.*

The schedules with $r_j = O(j)$ include, as far as we can tell, the cascades of Section 10.3. A positive answer would exclude every transposer whose motion looks like a cascade at all scales; a negative answer would be a transposer of that shape, faster than $O(L \log d)$.

Problem 11.2 (From the simulator back to the machine)**.** *Let M transpose in linear time, and let N_s be its bit simulator. Does the conclusion of Theorem 5.5 hold in M's own dependency DAG, on a set of traces followed by all but an exponentially small fraction of the inputs? A positive answer would remove the word "oblivious" from Theorem X. The question is whether a macroscopic crossing in the dependency DAG of N_s, which rewrites windows of $3s$ cells in each round, forces a macroscopic crossing of the heads of M.*

Problem 11.3 (Constants)**.** *Does Theorem 5.5 hold with a threshold $\theta \geq c_1/C$ and an expansion constant c_2, where $c_1, c_2 > 0$ depend polynomially on $1/k$, $1/w_{\mathrm{val}}$ and α? Numerically, the threshold of Theorem D appears not to be a universal constant (Remark 5.8), but nothing we know rules out thresholds of order $1/C$.*

Problem 11.4 (The cascades)**.** *For each of the five cascade schedules of Section 10.3, decide whether some machine with that oblivious trace transposes. Each is a linear-time schedule that met every necessary condition of this paper in our tests (Remark 10.1). A transposer with one of these traces would prove (T); an exclusion of one of them would, as far as our tests indicate, need an argument that uses more than Theorems D, X, S and Lemma R.*

A linear-time transposer, if one exists, must solve a coding problem on its own schedule: codes of rate approaching one on the cone relations of macroscopic queue-like reuse [6, Corollary O]. The fifteenth paper shows that binary storage codes of rate approaching one exist on suitable separated relations [14]; realizing such codes on the relations of a schedule that meets conditions (i)–(iv) of Section 10 would be a construction. In the other direction, showing that the cone relations of every such schedule, after a cheap genie, have a fixed rate deficit would exclude oblivious linear-time transposers (Theorem 3.4; [6, Corollary O]); deciding (T) itself needs the uniform native form of Theorem 9.5.

A Catalogue of excluded schedule families

This appendix lists the explicit oblivious linear-time schedule families excluded by the results of this paper. A schedule family prescribes, for each size d in its domain, the motion of the input head and of the work heads and the append steps; a family is *excluded* if for every $C \geq 1$ and every value width w there is d_0 such that, at every size $d \geq d_0$ in its domain, no machine of value width at most w whose trace is the prescribed one transposes within $C(L + 1)$ steps.

Status. For rows 1–20 the excluding theorem is formal; that the schedule belongs to the class is checked by inspection of the schedule, and for the few schedules defined only for even d or for powers of two (rows 3, 5 and 6) the class theorem is applied along those sizes, which the formal class predicates, quantified over all large d, do not literally cover. For rows 21–38 the certificate is built from formal lemmas of Section 3, but its verification on the schedule is a written argument; each is stated below as a proof sketch only. *None of these exclusions is a numbered result of this paper.*

Crossings (numerical). The last column records whether Theorem 5.5 alone could exclude the family. "Yes" means that the schedule's work heads stay within $O(b)$ cells, so it has no crossing revisits of linear excursion at all. The other entries are numerical, from exact computations at

$d \le 128$: "no" when the minimum set separating a large output set from the crossing steps of revisits of excursion more than $2W$ had at least 0.25 times as many elements as the output set; "small" when that ratio was at most 0.07, which suggests, as a heuristic, that Theorem 5.5 applies; and "—" when nothing was measured.

	Schedule	k	Excluded by	Crossings
1	Direct copy: output bit ℓ emitted while reading input bit ℓ; work tape idle (computes the identity)	1	Corollary 4.2, Theorem 9.2	yes
2	Persistent triangles: output (j, i, s) emitted while reading input $(j, -j - i \bmod d, s)$; work head within b cells	1	Corollary 4.2, Theorem 9.4	yes
3	One-pass shear with row rotations, two banks and a row buffer; seven monotone phases (d even)	3	Theorems 9.2 and 9.4	no
4	Ruler schedule: two banks moving one row per episode, a third head making excursions of 2-adic lengths inside one row	3	Theorem 9.4	—
5	Row expander: rows mixed in a two-head deep-nesting workspace of width W, two banks, four monotone phases (d even)	4	Theorem 9.2	—
6	Fan-out cascade fed by fresh input slices at each scale ($d = 2^m$)	2	Theorem 9.4 (suffix form)	—
7	Unit-slope passes: load, reflection, shifted sweep, emit; at most three constant-velocity segments per pass	2	Theorem 9.4	small
8	Bounce: zig-zag input, one tape sweeping $[0, L)$, nested excursions of height at most $L/(4d)$ on the other	2	Theorem 9.4, Corollary 4.2	small
9	Sublinear stack: blocks of $L/(4d)$ outputs pushed and popped on fresh cells while the other tape sweeps	2	Theorem 9.4, Corollary 4.2	—
10–11	Free pseudo-random Dyck walk of length L on one tape while the other sweeps the tape six times (two variants)	2	Theorem 9.4, Corollary 4.2	—
12	One push and one pop of height $2L$ on a stack tape while the other tape bounces	2	Theorem 9.4	—
13–16	One free head with all other excursions $O(L/\log d)$: slot cascades and skewed wavefronts (four variants)	2	Corollary 4.2	small (16)
17–20	No free head, owner load at most 4, family heights in at most 4 scale windows: multi-scale push–pop, two-excursion hub, LIFO Dyck stack, torus shifts	2	Theorem 9.6	—
21–23	One push–pop excursion on a stack tape, with a free tape running a multi-scale cascade or bouncing in a hub (three variants)	2	slab-chunk genie on a two-scale grid	—
24	Multi-scale Dyck stack: d sawtooth waves on a stack tape while the free tape runs a cascade and bounces in a hub ($d = 2^\kappa$)	2	block cut on the hub	no
25–26	Two push–pop excursions on a stack tape, with a mandatory cut at the second excursion (two variants)	2	block cut	—

	Schedule	k	Excluded by	Crossings
27	Two-exception inverse-square geometry, swept variant [8] ($d \geq 64$)	10	closed post-region, $O(d^{1/3})$ boundary steps	no
28–29	Two-exception inverse-square geometries [8] ($d \geq 64$)	9, 11	window lemma after two cut steps	small
30	Mirrored dyadic fold cascade	2	window lemma (Lemma 3.9), uncut	small
31	Unit-slope dyadic cascade ($d = 2^m$)	2	sub-row window, uncut	—
32	Dyadic fold cascade with row blocks of doubling width	2	window lemma, uncut	no
33–34	Two stack excursions with the free tape in a hub of width $L/2^s$ (s fixed) or d	2	block cut; space–time potential	—
35	Daisy-chained stack copies with golden-ratio shifts	2	piecewise-isometric windows	—
36	LIFO fan-out cascades of $\log d$ levels (two variants; $d = 2^\kappa$)	2	laminar cut of $2^{1-p}L + 3$ steps, $p = O(\log w_{\mathrm{val}})$	no
37	Equal-volume tapped cascade with $\log d$ meso-scale stages	2	meso-hub cut of $O(L/\log d)$ steps	no
38	Two-band meso cascade	2	closed post-region, $O(L/\log d)$ boundary steps	—

Rows 1–20. For bounded macro-phases (Theorems 9.2 and 9.4) the crux rate is $1/(4(|B|+1)^{M_{\mathrm{ph}}-1})$ with the bank B and the number of phases read off the schedule: for example, row 3 has banks $\{0,1\}$ and seven phases, row 5 banks $\{0,1\}$ and four phases. Rows 13–16 satisfy the hypotheses of Corollary 4.2 with the free head named and all other excursions $o(L)$; rows 17–20 those of Theorem 9.6 with owner load and window counts read off the schedule. For row 1, Theorem 9.2 gives crux rate $\frac{1}{4}$; the true rate is at most $\frac{1}{2}$, because pairing q with $\mathrm{tpos}(q)$ gives an exclusive-or code of $L/2$ bits.

Rows 21–38. Proof sketch only. Rows 21–23 are excluded by the frozen-trace count used in the proof of Theorem 4.1: on the pop leg of the stack, revisits that span the peak are order-reversing, and a two-scale slab grid with a greedy slicing gives a trapped chunk whose genie costs less than its output surplus. Rows 24–38 are excluded through Theorem 3.4 by an explicit certificate. Rows 24–26 and 33: the cut at a mandatory part of the stack motion makes every relevant output carry at most a blocks of potential sources and every input at most e blocks, and Lemma 3.8 gives an acyclic set of density $1/\binom{a+e}{a}$, uniformly in the depth. Rows 27–30 and 32: the window lemma (Lemmas 3.7 and 3.9) with potentials $\varphi(q) = \mathrm{tpos}(q)$ and $\psi(p) = L - 1 - p$ or a row-local variant, possibly after cutting a sublinear boundary. The geometries of rows 27–29 have hop counts of order L/i^2 in their i-th family; the gap-area memo of the ninth paper shows that its frozen-gap-area certificate cannot exclude this case [8], whereas the certificates here use cut steps and a window potential instead. Row 34 uses a space–time potential, row 35 a bound for piecewise-isometric motion with boundedly many segments, row 36 a bounded-level laminar cut, row 37 a meso-scale hub cut, and row 38 the formal closed-post-region criterion (`isAcyclicIn_ballAvoidSystem_of_closedPostRegion`).

As a numerical check on the method, the same classification was run on explicit transposers that run in time $\omega(L)$ (shuffle and butterfly transposers, a radix transposer on a stack, and one-tape

transposers); none of the exclusions was triggered by them.

B Index of Lean declarations

The formal proofs are in the Lean 4 library `NativeRelations` in `lean/` (Lean and Mathlib v4.34.1). It requires the libraries of the first, eighth, eleventh, fourteenth and fifteenth papers of the series by relative path. Its modules are, in the order of this paper: `DichromaticCrux.lean`, `MachineCrux.lean` and `EnvelopeCrux.lean` (Section 3); `OneFreeHeadSmallExcursion.lean` (Section 4); `ConjectureEBridge.lean` and `MacroscopicThreshold.lean` (Section 5); `MacroscopicSkeleton.lean` (Section 6); `TranspositionCut.lean` (Section 7); `SimulatorBridge.lean` (Section 8); `MacroBankCrux.lean`, `SuffixMassCrux.lean`, `JitterMacroBank.lean`, `AdaptiveSuffixJitterBM.lean`, `NativeGapCriterion.lean` and `ZeroFreeHeadQWindow.lean` (Section 9); and `ScaleProfile.lean` (Remark 8.3). The complete index of the declarations cited in this paper, with their files, will be generated from the sources.

C Building and verifying the formalization

The library builds with the Lean toolchain and the Mathlib revision pinned in `lean/lean-toolchain` and `lean/lake-manifest.json`, the same as the earlier papers of the series. Because it requires sibling libraries by relative path, the whole repository must be checked out. From `lean/`:

```
lake exe cache get
lake build
lake env lean PaperAxioms.lean
```

The last command prints, for every declaration cited in this paper, the axioms it depends on; each depends at most on `propext`, `Classical.choice` and `Quot.sound`. The library contains no `sorry`, no `native_decide`, no new axioms and no code that bypasses the kernel.

References

[1] Micah Adler, Nicholas J. A. Harvey, Kamal Jain, Robert D. Kleinberg, and April Rasala Lehman. On the capacity of information networks. In *Proceedings of the Seventeenth Annual ACM-SIAM Symposium on Discrete Algorithms (SODA 2006)*, pages 241–250, 2006.

[2] Alok Aggarwal and Jeffrey Scott Vitter. The input/output complexity of sorting and related problems. *Communications of the ACM*, 31(9):1116–1127, 1988. doi:10.1145/48529.48535.

[3] Ziv Bar-Yossef, Yitzhak Birk, T. S. Jayram, and Tomer Kol. Index coding with side information. *IEEE Transactions on Information Theory*, 57(3):1479–1494, 2011.

[4] Mark Braverman and Zhongtian He. The multiple unicast conjecture is false. arXiv:2610.10108, 2026. URL: https://arxiv.org/abs/2610.10108.

[5] Hans Brende and Claude Opus 5.5 (Anthropic). Bounded nesting and depth reduction for matrix transposition on tapes. Manuscript and Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Bounded-Nesting-and-Depth-Reduction-for-Matrix-Transposition-on-Tapes.

[6] Hans Brende and Claude Opus 5.5 (Anthropic). Bounded-stack DAGs reduce to constant depth. Manuscript and Lean formalization, repository HansBrende/papers, November 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Bounded-Stack-DAGs-Reduce-to-Constant-Depth.

[7] Hans Brende and Claude Opus 5.5 (Anthropic). Curved motion, depth reduction, and storage codes for matrix transposition on tapes. Manuscript and partial Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Curved-Motion-Depth-Reduction-and-Storage-Codes-for-Matrix-Transposition-on-Tapes.

[8] Hans Brende and Claude Opus 5.5 (Anthropic). Frozen gap areas and the sharpness of source-cover certificates. Memo to the ninth paper of the series, with Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Source-Covers-and-Barriers-to-Tape-Transposition/memo-gap-area.

[9] Hans Brende and Claude Opus 5.5 (Anthropic). Matrix transposition and the tape coding conjecture. Manuscript and Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Matrix-Transposition-and-the-Tape-Coding-Conjecture.

[10] Hans Brende and Claude Opus 5.5 (Anthropic). Merge depth, input steering, and the coding core of matrix transposition on tapes. Manuscript and Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Merge-Depth-Input-Steering-and-the-Coding-Core-of-Matrix-Transposition-on-Tapes.

[11] Hans Brende and Claude Opus 5.5 (Anthropic). Storage codes above half rate under paired-distance separation. Manuscript and Lean formalization, repository HansBrende/papers, November 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Storage-Codes-above-Half-Rate-under-Paired-Distance-Separation.

[12] Hans Brende and Gemini 4 Argon (Google DeepMind). Directed reciprocity and deep-nesting bounds for matrix transposition on tapes. Manuscript and Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Directed-Reciprocity-and-Deep-Nesting-Bounds-for-Matrix-Transposition-on-Tapes.

[13] Hans Brende and GPT-6 Astra (OpenAI). Adaptive pass bounds without a runtime assumption. Memo to the first paper of the series, with Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Matrix-Transposition-and-the-Tape-Coding-Conjecture/memo-passes.

[14] Hans Brende and GPT-6 Astra (OpenAI). High-rate separated storage codes and tape transposition. Manuscript and Lean formalization, repository HansBrende/papers, November 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/High-Rate-Separated-Storage-Codes-and-Tape-Transposition.

[15] Hans Brende and GPT-6 Astra (OpenAI). Source covers and barriers to tape transposition. Manuscript and Lean formalization, repository HansBrende/papers, October 2026. Pinned revision d5576b2391fb81e782540339776318a2abd065a2. URL: https://github.com/HansBrende/papers/tree/d5576b2391fb81e782540339776318a2abd065a2/Source-Covers-and-Barriers-to-Tape-Transposition.

[16] Martin Dietzfelbinger and Wolfgang Maass. The complexity of matrix transposition on one-tape off-line Turing machines with output tape. *Theoretical Computer Science*, 108(2):271–290, 1993.

[17] Martin Dietzfelbinger, Wolfgang Maass, and Georg Schnitger. The complexity of matrix transposition on one-tape off-line Turing machines. *Theoretical Computer Science*, 82(1):113–129, 1991.

[18] Étienne Grandjean. A natural NP-complete problem with a nontrivial lower bound. *SIAM Journal on Computing*, 17(4):786–809, 1988. doi:10.1137/0217050.

[19] Martin Grohe, André Hernich, and Nicole Schweikardt. Lower bounds for processing data with few random accesses to external memory. *Journal of the ACM*, 56(3), 2009. Article 12.

[20] Martin Grohe and Nicole Schweikardt. Lower bounds for sorting with few random accesses to external memory. In *Proceedings of the Twenty-fourth ACM SIGMOD-SIGACT-SIGART Symposium on Principles of Database Systems (PODS 2005)*. ACM, 2005. doi:10.1145/1065167.1065197.

[21] David Harvey and Joris van der Hoeven. Integer multiplication is at least as hard as matrix transposition. arXiv:2503.22848; also in Proc. FOCS 2025, pp. 1366–1381 (numbering as in arXiv), 2025. URL: https://arxiv.org/abs/2503.22848.

[22] F. C. Hennie. One-tape, off-line Turing machine computations. *Information and Control*, 8(6):553–578, 1965.

[23] John Hopcroft, Wolfgang Paul, and Leslie Valiant. On time versus space. *Journal of the ACM*, 24(2):332–337, 1977. doi:10.1145/322003.322015.

[24] Walter W. Kirchherr. Transposition of an $l \times l$ matrix requires $\Omega(\log l)$ reversals on conservative Turing machines. *Information Processing Letters*, 28(2):55–59, 1988. doi:10.1016/0020-0190(88)90163-9.

[25] Wolfgang Maass. Combinatorial lower bound arguments for deterministic and nondeterministic Turing machines. *Transactions of the American Mathematical Society*, 292(2):675–693, 1985.

[26] OpenAI. Integer multiplication below $n \log n$. GitHub release, repository `openai/math`, family 109, September 2026. Preprint dated 23 September 2026. URL: https://github.com/openai/math/blob/main/preprints/Integer-multiplication-below-n-log-n-September-23-2026/paper.pdf.

[27] Wolfgang J. Paul, Nicholas Pippenger, Endre Szemerédi, and William T. Trotter. On determinism versus non-determinism and related problems. In *24th Annual Symposium on Foundations of Computer Science (FOCS 1983)*, pages 429–438, 1983.

[28] Nicholas Pippenger and Michael J. Fischer. Relations among complexity measures. *Journal of the ACM*, 26(2):361–381, 1979.

[29] Hans-Jörg Stoß. Rangierkomplexität von Permutationen. *Acta Informatica*, 2:80–96, 1973.

[30] Xindan Zhang, Baochun Li, and Zongpeng Li. Network coding can beat routing in undirected multiple-unicast networks. arXiv:2610.09367, 2026. URL: https://arxiv.org/abs/2610.09367.